

# Embedded Software

---

Introduction to operating systems

# Agenda

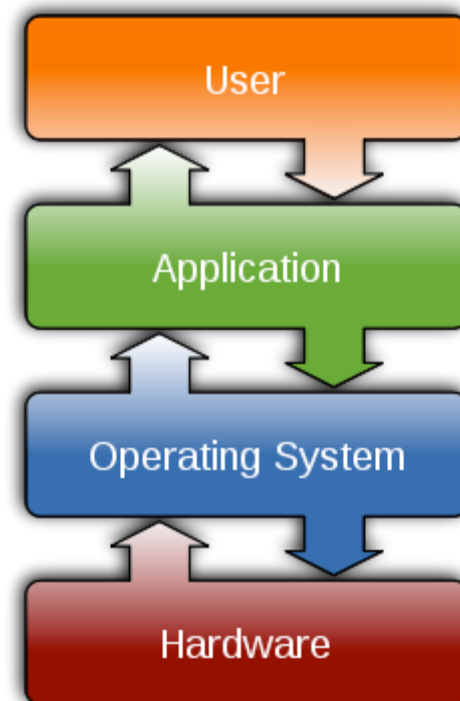
---

- What is an operating system?
- OS structure
  - ▶ Process management
  - ▶ Memory and storage mgmt.
  - ▶ I/O subsystem
- Operating systems - Real-Time OS's

# What is an operating system?

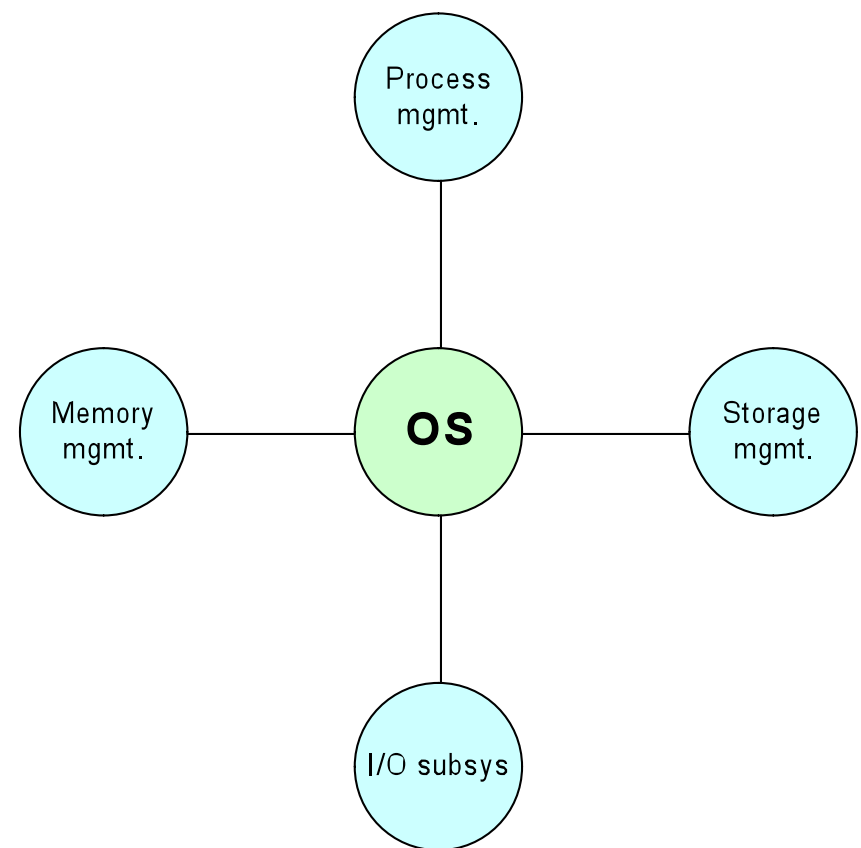
---

- What is an operating system?
  - ▶ Wikipedia: *“An operating system (OS) is software (...) that manages computer hardware resources and provides common services for efficient execution of various application software.”*



# OS structure

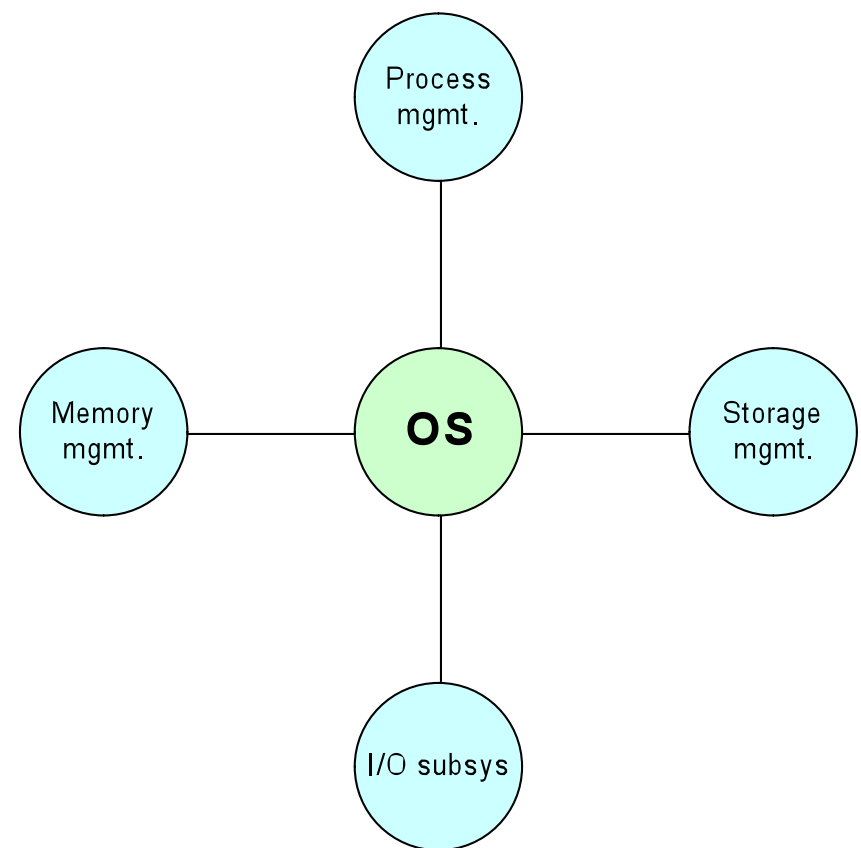
---



# OS structure

---


- Many computer types – many OS designs
  - ▶ Mainframe OSs are optimized for HW utilization
  - ▶ Desktop OSs are optimized for generality
  - ▶ Embedded OSs are optimized for efficiency, size, safety, speed, low power
  - ▶ ...



# OS structure

---

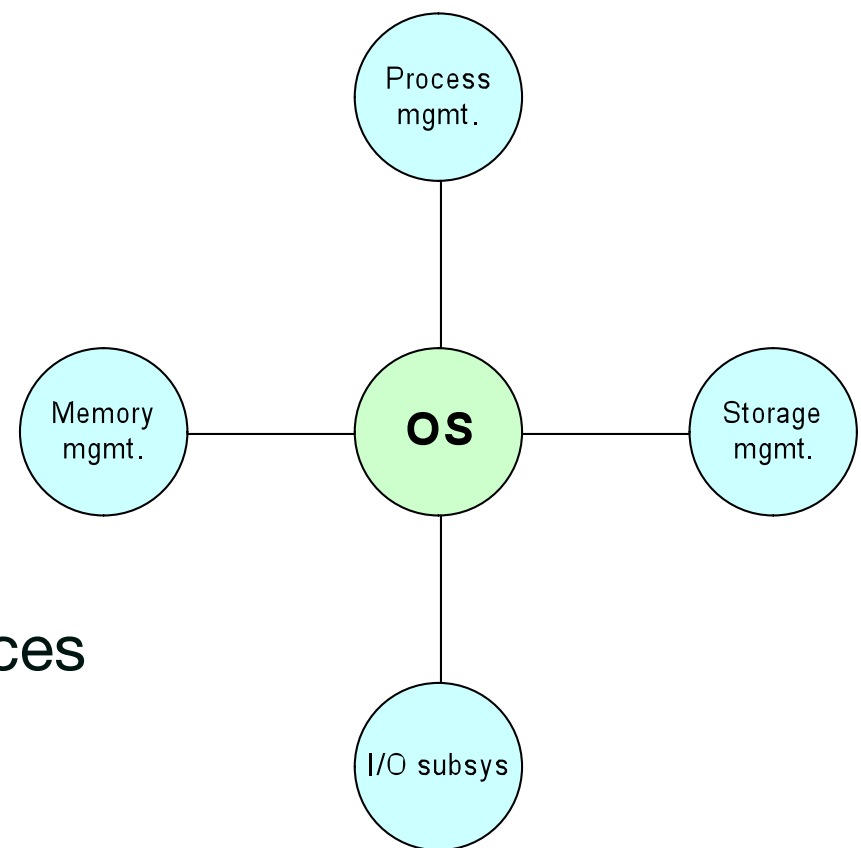
- Many computer types – many OS designs
  - ▶ Mainframe OSs are optimized for HW utilization
  - ▶ Desktop OSs are optimized for generality
  - ▶ Embedded OSs are optimized for efficiency, size, safety, speed, low power
  - ▶ ...
- Some commonalities, though:
  - ▶ Process management – handles multiprogramming and keep the CPU busy
  - ▶ Memory management – (de)allocation and process swapping
  - ▶ Storage management – persistent storage and cache
  - ▶ I/O subsystem management – manage I/O devices



# OS structure

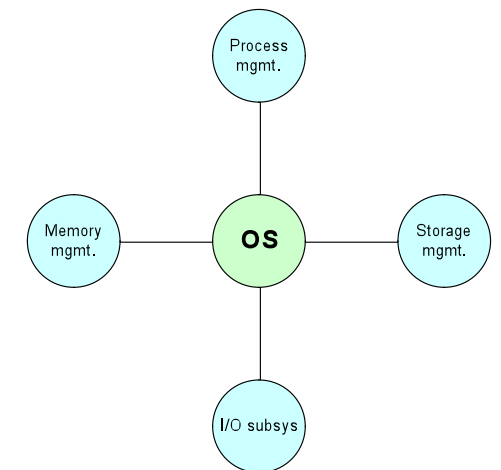
---

- Many computer types – many OS designs
  - ▶ Mainframe OSs are optimized for HW utilization
  - ▶ Desktop OSs are optimized for generality
  - ▶ Embedded OSs are optimized for efficiency, size, safety, speed, low power
  - ▶ ...
- Some commonalities, though:
  - ▶ Process management – handles multiprogramming and keep the CPU busy
  - ▶ Memory management – (de)allocation and process swapping
  - ▶ Storage management – persistent storage and cache
  - ▶ I/O subsystem management – manage I/O devices
- Let's take a look at process management



# OS structure - Process management

---






# OS structure - Process management

---

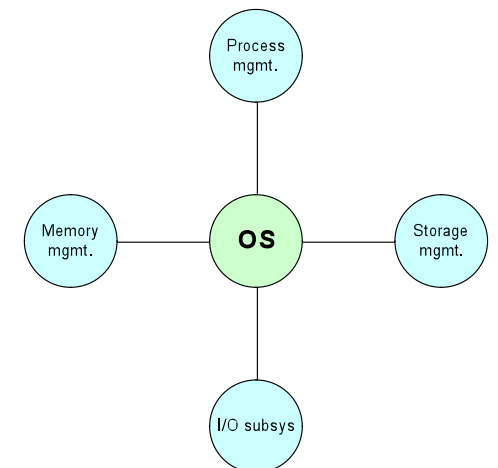
- ***What is a process? Is it a program?***
  - ▶ No - a process is a program in execution



# OS structure - Process management

---

- ***What is a process? Is it a program?***
  - ▶ No - a process is a program in execution



# OS structure - Process management

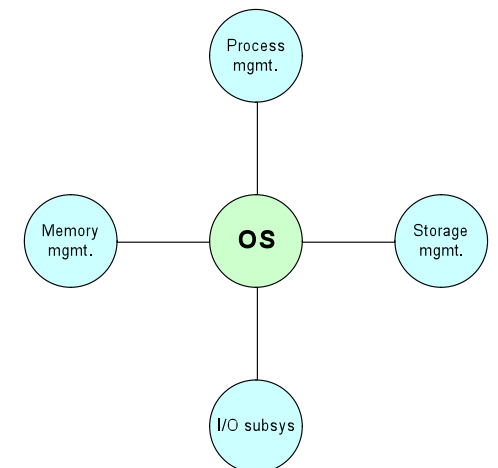
---

- ***What is a process? Is it a program?***

- ▶ No - a process is a program in execution

- ***How many processes can run at a time?***

- ▶ There can be many processes that want to run, but only one per CPU that actually runs



# OS structure - Process management

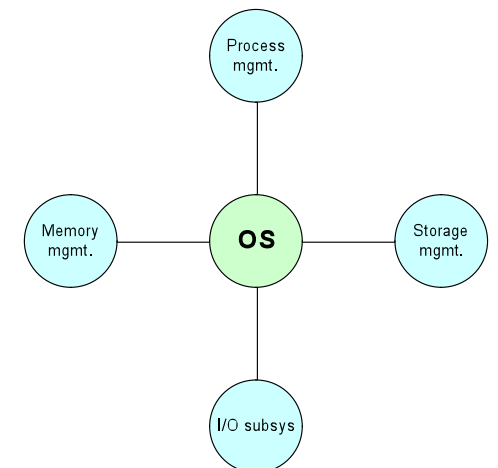
---

- ***What is a process? Is it a program?***

- ▶ No - a process is a program in execution

- ***How many processes can run at a time?***

- ▶ There can be many processes that want to run, but only one per CPU that actually runs



# OS structure - Process management

---

- ***What is a process? Is it a program?***

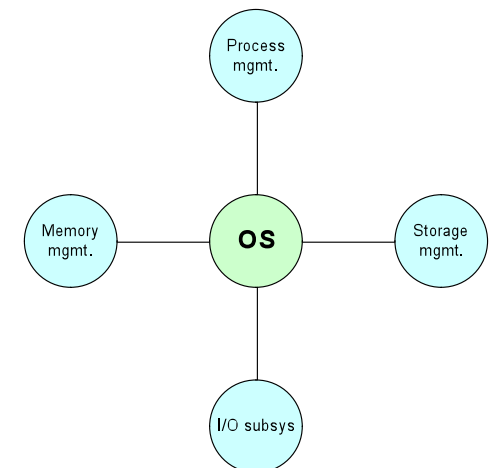
- ▶ No - a process is a program in execution

- ***How many processes can run at a time?***

- ▶ There can be many processes that want to run, but only one per CPU that actually runs

- The OS manages processes

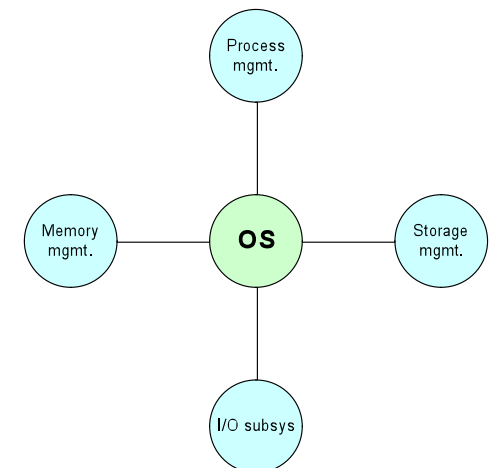
- ▶ Creates, deletes, and allocates resources for them
- ▶ Swaps them in and out of memory
- ▶ Suspends and resumes them
- ▶ Provides mechanisms for synchronization and communication between processes



# Process management – why?

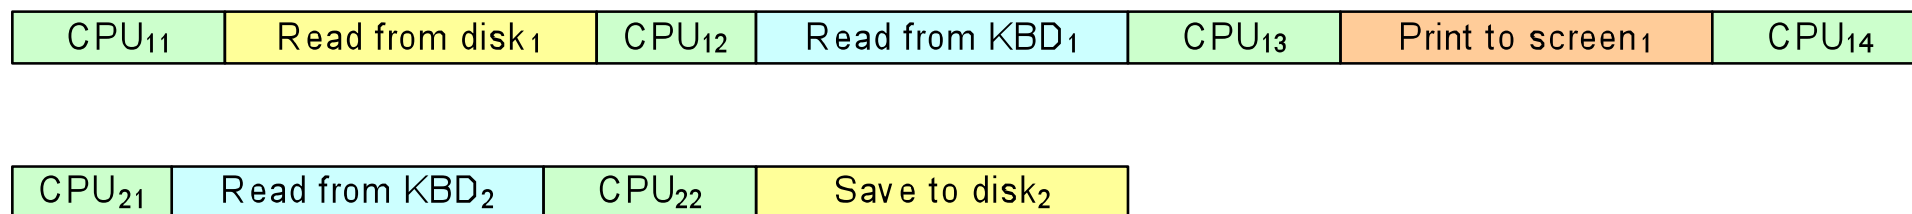
---

- Processes either compute or perform device I/O
- ***What does a process do while it performs I/O?***
  - ▶ It must wait for I/O to complete before it can resume
- ***What should the system do meanwhile?***
  - ▶ Without process management: CPU idles
  - ▶ With process management: Switch to another ready process




# Process management - example

- Consider two tasks:

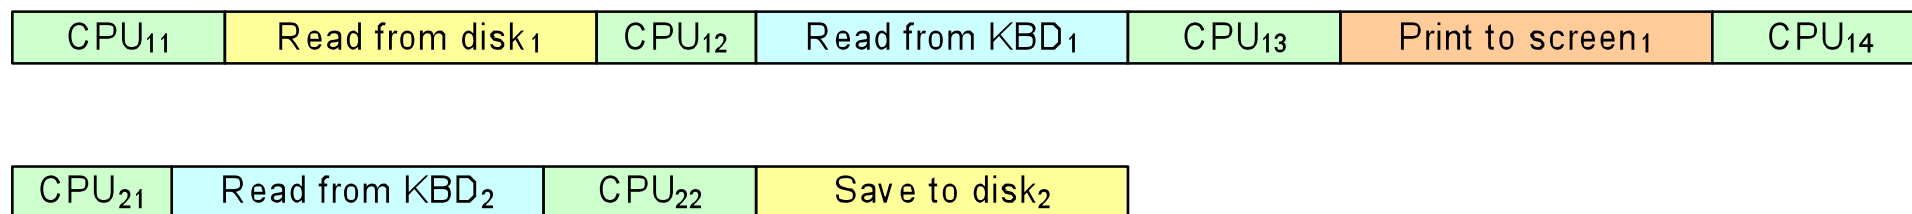


- Scheduling without resource management (batch processing)?***

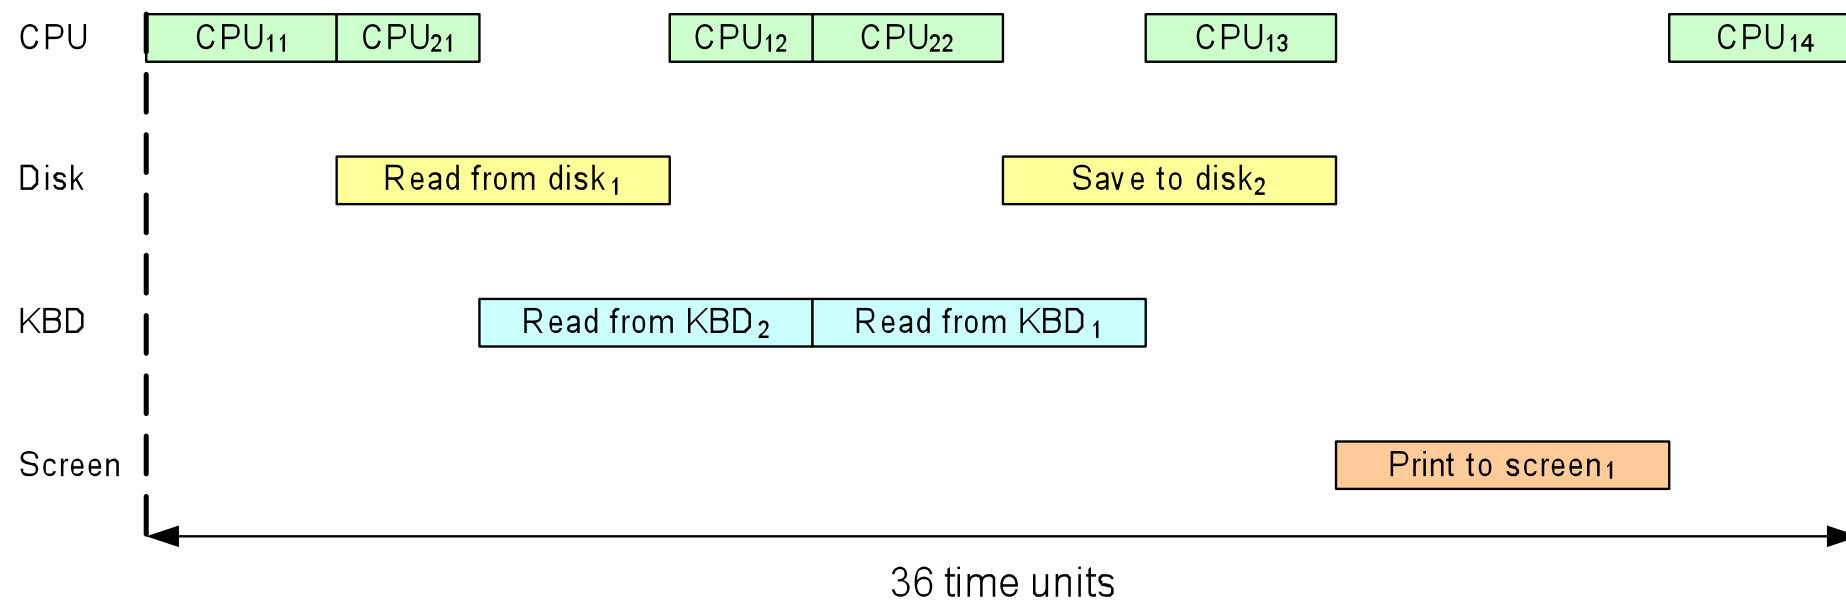


# Process management - example

- Consider two tasks:



- Scheduling with resource management?***





# Process management - protection

---

# Process management - protection

---

- Consider an “evil” process – what damage could it do?
  - ▶ Destroy, eavesdrop on, change other processes
  - ▶ Destroy OS
  - ▶ Destroy files and HW

# Process management - protection

---

- Consider an “evil” process – what damage could it do?
  - ▶ Destroy, eavesdrop on, change other processes
  - ▶ Destroy OS
  - ▶ Destroy files and HW

# Process management - protection

---

- Consider an “evil” process – what damage could it do?
  - ▶ Destroy, eavesdrop on, change other processes
  - ▶ Destroy OS
  - ▶ Destroy files and HW
- The OS guards against this using dual-mode operation (MODE bit in CPU)
  - ▶ Applications run in user mode (AKA restricted mode)
  - ▶ The OS kernel runs in kernel mode (AKA protected, privileged, supervisor mode)

# Process management - protection

---

- Consider an “evil” process – what damage could it do?
  - ▶ Destroy, eavesdrop on, change other processes
  - ▶ Destroy OS
  - ▶ Destroy files and HW
- The OS guards against this using dual-mode operation (MODE bit in CPU)
  - ▶ Applications run in user mode (AKA restricted mode)
  - ▶ The OS kernel runs in kernel mode (AKA protected, privileged, supervisor mode)

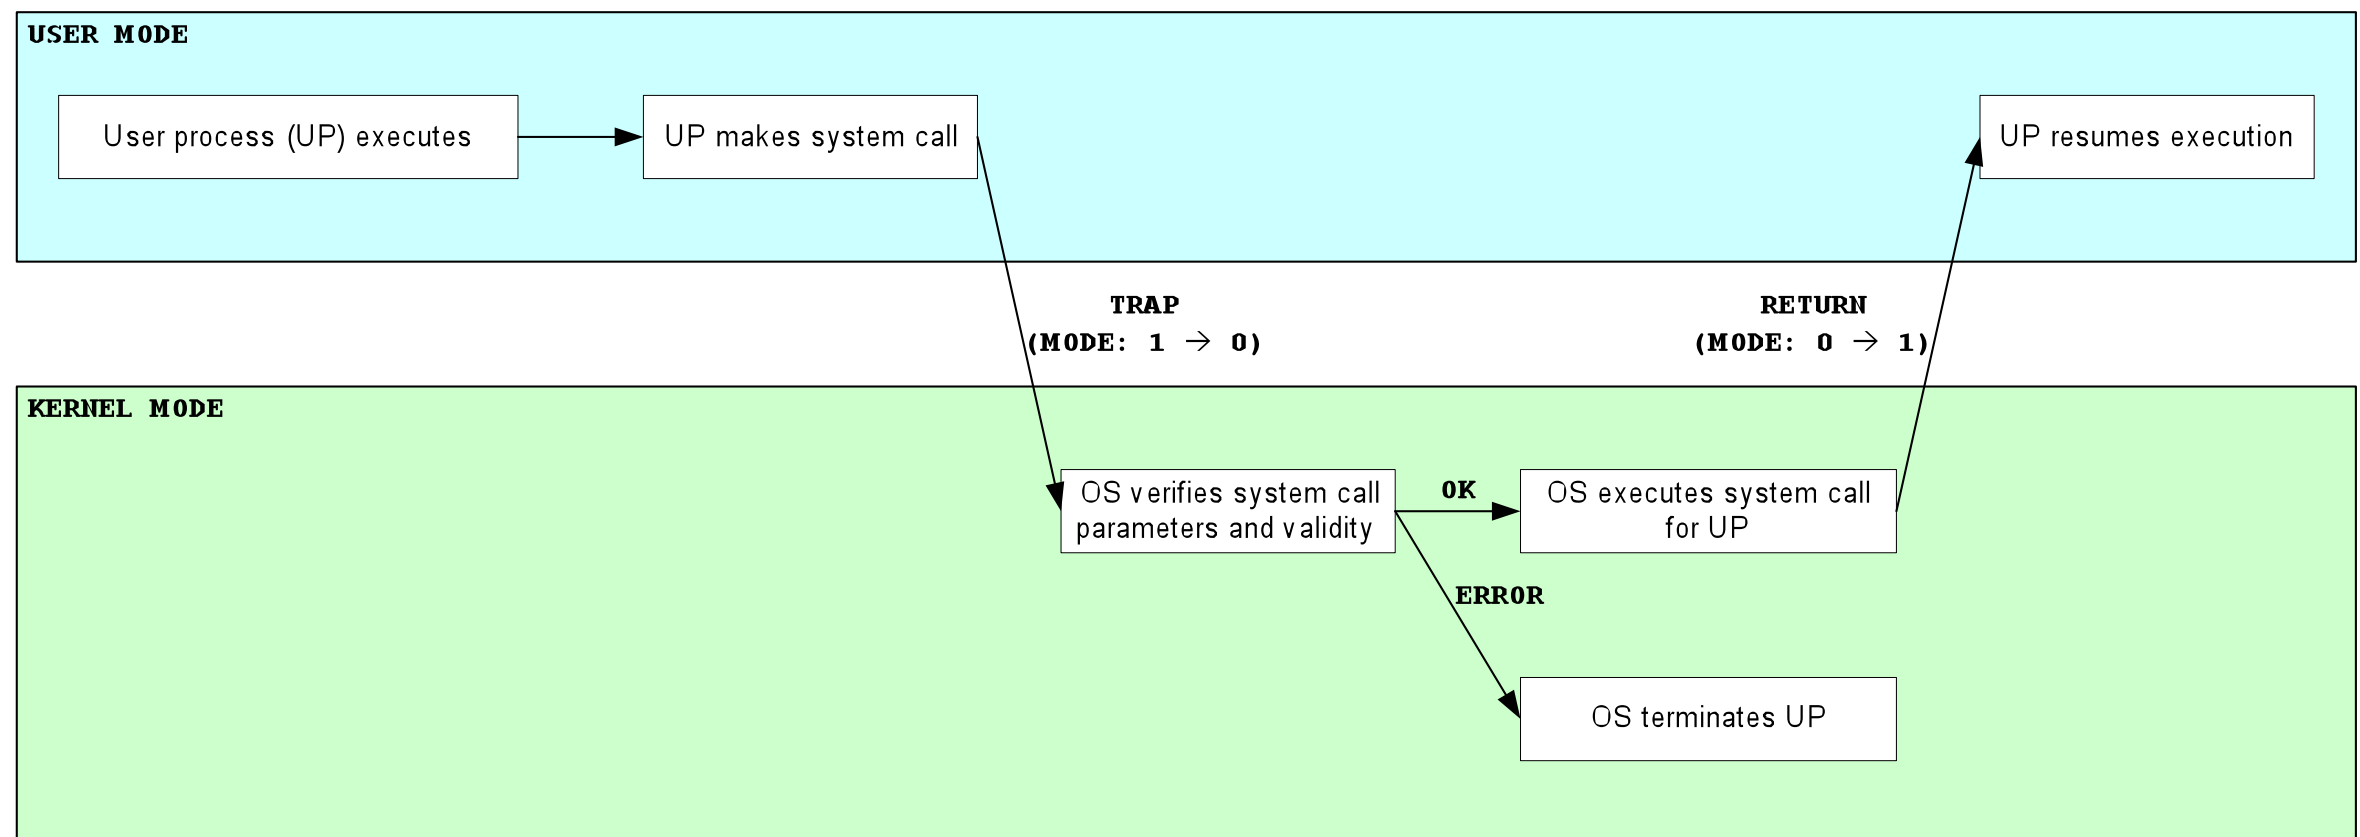
# Process management - protection

---

- Consider an “evil” process – what damage could it do?
  - ▶ Destroy, eavesdrop on, change other processes
  - ▶ Destroy OS
  - ▶ Destroy files and HW
- The OS guards against this using dual-mode operation (MODE bit in CPU)
  - ▶ Applications run in user mode (AKA restricted mode)
  - ▶ The OS kernel runs in kernel mode (AKA protected, privileged, supervisor mode)
- Potentially dangerous operations (I/O, IPC, ...) can only be done via privileged instructions
  - ▶ Restricted instructions – user and kernel mode
  - ▶ Privileged instructions – kernel mode only

# Dual-mode operation – system calls

- When processes need to perform I/O, it does so via the OS via well-defined **system calls** (version 2.6.35: 337 different syscalls)
- The OS (which is in kernel mode) verifies the system call and its parameters



# Dual-mode operation – system calls

---

- How often are system calls made?



# Dual-mode operation – system calls

---

- How often are system calls made?

```
$ ./hello  
Hello World!
```

```
$
```

# Dual-mode operation – system calls

---

- How often are system calls made?

```
$ ./hello  
Hello World!  
$ strace ./hello
```

```
$
```

# Dual-mode operation – system calls

- How often are system calls made?

```
$ ./hello
Hello World!
$ strace ./hello
execve("././hello", ["../hello"], [/* 46 vars */]) = 0
brk(0) = 0x8568000
access("/etc/ld.so.nohwcap", F_OK) = -1 ENOENT (No such file or
directory)
mmap2(NULL, 8192, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) = 0xb77ef000
access("/etc/ld.so.preload", R_OK) = -1 ENOENT (No such file or directory)
open("/etc/ld.so.cache", O_RDONLY) = 3
fstat64(3, {st_mode=S_IFREG|0644, st_size=57413, ...}) = 0
mmap2(NULL, 57413, PROT_READ, MAP_PRIVATE, 3, 0) = 0xb77ef000
close(3) = 0
access("/etc/ld.so.nohwcap", F_OK) = -1 ENOENT (No such file or directory)
open("/usr/lib/libstdc++.so.6", O_RDONLY) = 3
read(3, "\177ELF\1\1\1\0\0\0\0\0\0\0\3\0\3\0\1\0\0\0\0\0\4\0\0004\0\0\0"... 512) = 512
fstat64(3, {st_mode=S_IFREG|0644, st_size=75989, ...}) = 0
mmap2(NULL, 180428, PROT_READ|PROT_EXEC, MAP_PRIVATE|MAP_DENYWRITE, 3, 0) = 0xbbf000
mprotect(0xbbf000, 4096, PROT_NONE) = 0
mmap2(0xbbf000, 26400, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_FIXED|MAP_DENYWRITE, 3, 0xbbf) = 0xbbf000
mmap2(0xbbf000, 25484, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_FIXED|MAP_ANONYMOUS, -1, 0) = 0xbbf000
close(3) = 0
access("/etc/ld.so.nohwcap", F_OK) = -1 ENOENT (No such file or directory)
open("/lib/tls/i686/cmov/libc.so.6", O_RDONLY) = 3
read(3, "\177ELF\1\1\1\0\0\0\0\0\0\0\3\0\3\0\1\0\0\0\0\0\4\0\0004\0\0\0"... 512) = 512
fstat64(3, {st_mode=S_IFREG|0644, st_size=143392, ...}) = 0
mmap2(NULL, 151680, PROT_READ|PROT_EXEC, MAP_PRIVATE|MAP_DENYWRITE, 3, 0) = 0xb7d4000
mmap2(0xb7d4000, 8192, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_FIXED|MAP_DENYWRITE, 3, 0xb7d) = 0xb7d4000
close(3) = 0
access("/etc/ld.so.nohwcap", F_OK) = -1 ENOENT (No such file or directory)
open("/lib/tls/i686/cmov/libc.so.6", O_RDONLY) = 3
read(3, "\177ELF\1\1\1\0\0\0\0\0\0\0\3\0\3\0\1\0\0\0\0\0\4\0\0004\0\0\0"... 512) = 512
fstat64(3, {st_mode=S_IFREG|0644, st_size=120368, ...}) = 0
mmap2(NULL, 4096, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) = 0xb77ef000
mmap2(NULL, 123432, PROT_READ|PROT_EXEC, MAP_PRIVATE|MAP_DENYWRITE, 3, 0) = 0xb7f000
mmap2(0xb7f000, 8192, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_FIXED|MAP_DENYWRITE, 3, 0xb7f) = 0xb7f000
close(3) = 0
access("/etc/ld.so.nohwcap", F_OK) = -1 ENOENT (No such file or directory)
open("/lib/tls/i686/cmov/libc.so.6", O_RDONLY) = 3
read(3, "\177ELF\1\1\1\0\0\0\0\0\0\0\3\0\3\0\1\0\0\0\0\0\4\0\0004\0\0\0"... 512) = 512
fstat64(3, {st_mode=S_IFREG|0755, st_size=140580, ...}) = 0
mmap2(NULL, 141552, PROT_READ|PROT_EXEC, MAP_PRIVATE|MAP_DENYWRITE, 3, 0) = 0x110000
mprotect(0x110000, 4096, PROT_NONE) = 0
mmap2(0x110000, 12288, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_FIXED|MAP_DENYWRITE, 3, 0x110) = 0x110000
mmap2(0x110000, 18664, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_FIXED|MAP_ANONYMOUS, -1, 0) = 0x110000
close(3) = 0
mmap2(NULL, 4096, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) = 0xb77ef000
set_thread_area({entry_number:-1 -> 0, base_addr:0xb77f66d0, limit:140575, seg_32bit:1, contents:0, read_exec_only:0, limit_in_pages:1, seg_not_present:0, useable:1}) = 0
mprotect(0x110000, 8192, PROT_READ) = 0
mprotect(0xb7ae000, 4096, PROT_READ) = 0
mprotect(0xb7ae000, 4096, PROT_READ) = 0
mprotect(0xb7ae000, 16384, PROT_READ) = 0
mprotect(0xb7ae000, 4096, PROT_READ) = 0
mprotect(0xb7ae000, 4096, PROT_READ) = 0
munmap(0xb77ef000, 57413) = 0
fstat64(1, {st_mode=S_IFCHR|0620, st_rdev=unknown, ...}) = 0
mmap2(NULL, 4096, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) = 0xb77ef000
write(1, "Hello world!\n", 13Hello world!
) = 13
exit_group(0) =
$
```

# Dual-mode operation – system calls



- How often are system calls made?

# 56 system calls

```
$ ./hello
Hello World!
$ strace ./hello
execve("././hello", ["../hello"], [/* 46 vars */]) = 0
brk(0) = 0x8568000
access("/etc/ld.so.nohwcap", F_OK) = -1 ENOENT (No such file or
directory)
mmap2(NULL, 8192, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) = 0xb77ef000
access("/etc/ld.so.preload", R_OK) = -1 ENOENT (No such file or directory)
open("/etc/ld.so.cache", O_RDONLY) = 3
fstat64(3, {st_mode=S_IFREG|0644, st_size=57413, ...}) = 0
mmap2(NULL, 57413, PROT_READ, MAP_PRIVATE, 3, 0) = 0xb77ef000
close(3) = 0
access("/etc/ld.so.nohwcap", F_OK) = -1 ENOENT (No such file or directory)
open("/usr/lib/libstdc++.so.6", O_RDONLY) = 3
read(3, "\177ELF\1\1\1\0\0\0\0\0\0\0\3\0\3\0\1\0\0\0\0\0\4\0\004\0\0\0.... 512) = 512
fstat64(3, {st_mode=S_IFREG|0644, st_size=75989, ...}) = 0
mmap2(NULL, 180428, PROT_READ|PROT_EXEC, MAP_PRIVATE|MAP_DENYWRITE, 3, 0) = 0xbbf000
mprotect(0xbbf000, 4096, PROT_NONE) = 0
mmap2(0xbbf000, 2640, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_FIXED|MAP_DENYWRITE, 3, 0xbbf) = 0xbbf000
mmap2(0xbbf000, 25484, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_FIXED|MAP_ANONYMOUS, -1, 0) = 0xbbf000
close(3) = 0
access("/etc/ld.so.nohwcap", F_OK) = -1 ENOENT (No such file or directory)
open("/lib/ls/1685/cnv/libc.so.6", O_RDONLY) = 3
read(3, "\177ELF\1\1\1\0\0\0\0\0\0\0\3\0\3\0\1\0\0\0\0\0\4\0\004\0\0\0.... 512) = 512
fstat64(3, {st_mode=S_IFREG|0644, st_size=143392, ...}) = 0
mmap2(NULL, 151680, PROT_READ|PROT_EXEC, MAP_PRIVATE|MAP_DENYWRITE, 3, 0) = 0xb7d4000
mmap2(0xb7d4000, 8192, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_FIXED|MAP_DENYWRITE, 3, 0x23) = 0xb7d4000
close(3) = 0
access("/etc/ld.so.nohwcap", F_OK) = -1 ENOENT (No such file or directory)
open("/lib/ls/libgcc_s.so.1", O_RDONLY) = 3
read(3, "\177ELF\1\1\1\0\0\0\0\0\0\0\3\0\3\0\1\0\0\0\0\0\4\0\004\0\0\0.... 512) = 512
fstat64(3, {st_mode=S_IFREG|0644, st_size=120368, ...}) = 0
mmap2(NULL, 4096, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) = 0xb77ef000
mmap2(NULL, 123432, PROT_READ|PROT_EXEC, MAP_PRIVATE|MAP_DENYWRITE, 3, 0) = 0xb7f1000
mmap2(0xb7f1000, 8192, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_FIXED|MAP_DENYWRITE, 3, 0xc1) = 0xb7f1000
close(3) = 0
access("/etc/ld.so.nohwcap", F_OK) = -1 ENOENT (No such file or directory)
open("/lib/ls/1685/cnv/libc.so.6", O_RDONLY) = 3
read(3, "\177ELF\1\1\1\0\0\0\0\0\0\0\3\0\3\0\1\0\0\0\0\0\4\0\004\0\0\0.... 512) = 512
fstat64(3, {st_mode=S_IFREG|0755, st_size=140580, ...}) = 0
mmap2(NULL, 141592, PROT_READ|PROT_EXEC, MAP_PRIVATE|MAP_DENYWRITE, 3, 0) = 0x110000
mprotect(0x110000, 4096, PROT_NONE) = 0
mmap2(0x110000, 12288, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_FIXED|MAP_DENYWRITE, 3, 0x153) = 0x264000
mmap2(0x264000, 18664, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_FIXED|MAP_ANONYMOUS, -1, 0) = 0x267000
close(3) = 0
mmap2(NULL, 4096, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) = 0xb77ef000
set_thread_area({entry_number:-1 -> 0, base_addr:0xb77f66d0, limit:140575, seg_3bit:1, contents:0, read_exec_only:0, limit_in_pages:1, seg_not_present:0, useable:1}) = 0
mprotect(0x264000, 8192, PROT_READ) = 0
mprotect(0xb7ae000, 4096, PROT_READ) = 0
mprotect(0xb7ae000, 4096, PROT_READ) = 0
mprotect(0xb7ae000, 16384, PROT_READ) = 0
mprotect(0xb7ae000, 4096, PROT_READ) = 0
mprotect(0xb7ae000, 4096, PROT_READ) = 0
munmap(0xb77ef000, 57413) = 0
fstat64(1, {st_mode=S_IFCHR|0620, st_rdev=unknown (135, 0), ...}) = 0
mmap2(NULL, 4096, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) = 0xb77ef000
write(1, "Hello world!\n", 13Hello world!
) = 13
exit_group(0) =
$
```

# OS Structure - Memory and storage mgmt.

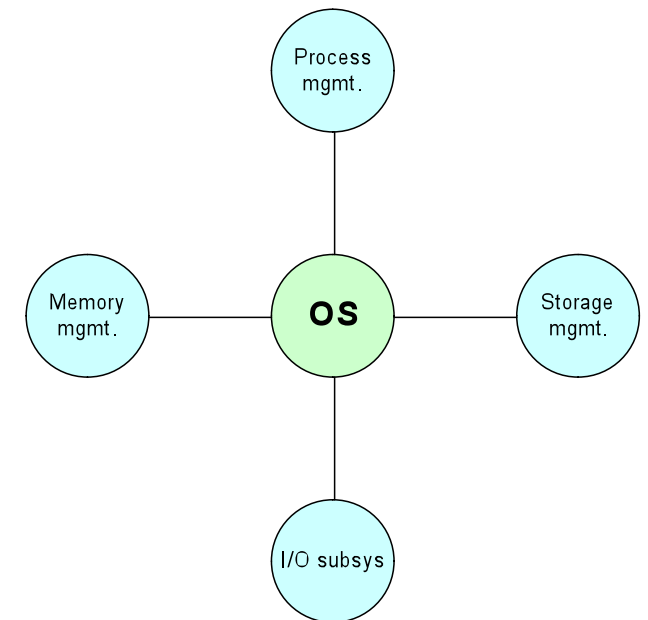
- Memory management
  - ▶ Keep track of several processes in memory at a time
  - ▶ Decide which (parts of) processes to move in and out of memory
    - ▶ Many different algorithms depending on hardware and the OS' nature
  - ▶ Allocate and deallocate memory as necessary
- Storage management
  - ▶ The primary storage is never big enough to accommodate all needs
  - ▶ A hierarchy of memory:
    - ▶ Size? Price (per MB)? Capacity? Bandwidth?
  - ▶ Move data in/out of hierarchy



# OS Structure - I/O subsystem

---

- The I/O subsystem hides the oddities of individual I/O devices
- Instead, it provides a uniform interface (in Linux: a file)
  - ▶ The file is I/F to a device driver
  - ▶ The device driver knows how to operate the device



# Operating systems - Real-Time OS's

---

- Real-Time Operating Systems (RTOSs) are OSs intended for RT systems (!)
- ***Some key properties?***
  - ▶ Minimum interrupt latency
  - ▶ Minimum task switching latency
    - ▶ Includes known worst case latency (must be small)
  - ▶ Static task priorities
- The programmer (you!) is responsible for correct priority assignment