

Embedded Software


Parallel programs, processes and threads

Agenda

- The problem - A Case
- A Solution - Parallelism
- Processes and Threads in Linux
- Advantages & Disadvantages with multitasking

Case

- Consider a system that allows a user to enter a PIN.
 - ***How would you implement this?***

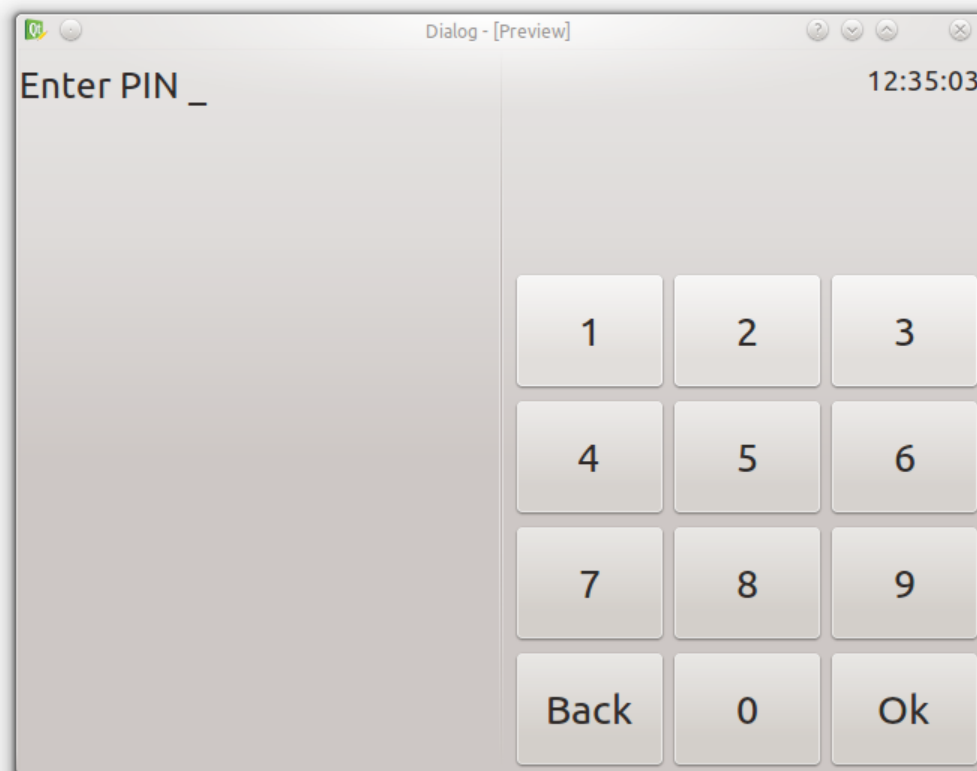


```
void main()
{
    print "ENTER PIN:"
    get key -> ch
    while(ch != "OK")
        input += ch
        compare(input, pin)
        ...
}
```

Case

- Now consider the same system, but now with a clock.

‣ ***How would you implement this?***




```
void main()  
{  
    ???  
}
```

- How would you update the clock while waiting for input?
- How would you capture key presses while updating the clock?

Case - Solution

- With multiple threads, this is trivial:



Handles user input as before

```
void userInput()  
{  
    print "ENTER PIN:"  
    get key -> ch  
    while(ch != "OK")  
        input += ch  
        compare(input, pin)  
    ...  
}
```

Handles clock updates

```
void updClock()  
{  
    while(true)  
    {  
        display current time  
        wait 1s  
    }  
}
```

Starts threads that run user input and clock update routines

```
int main()  
{  
    startThread(userInput)  
    startThread(updClock)  
}
```

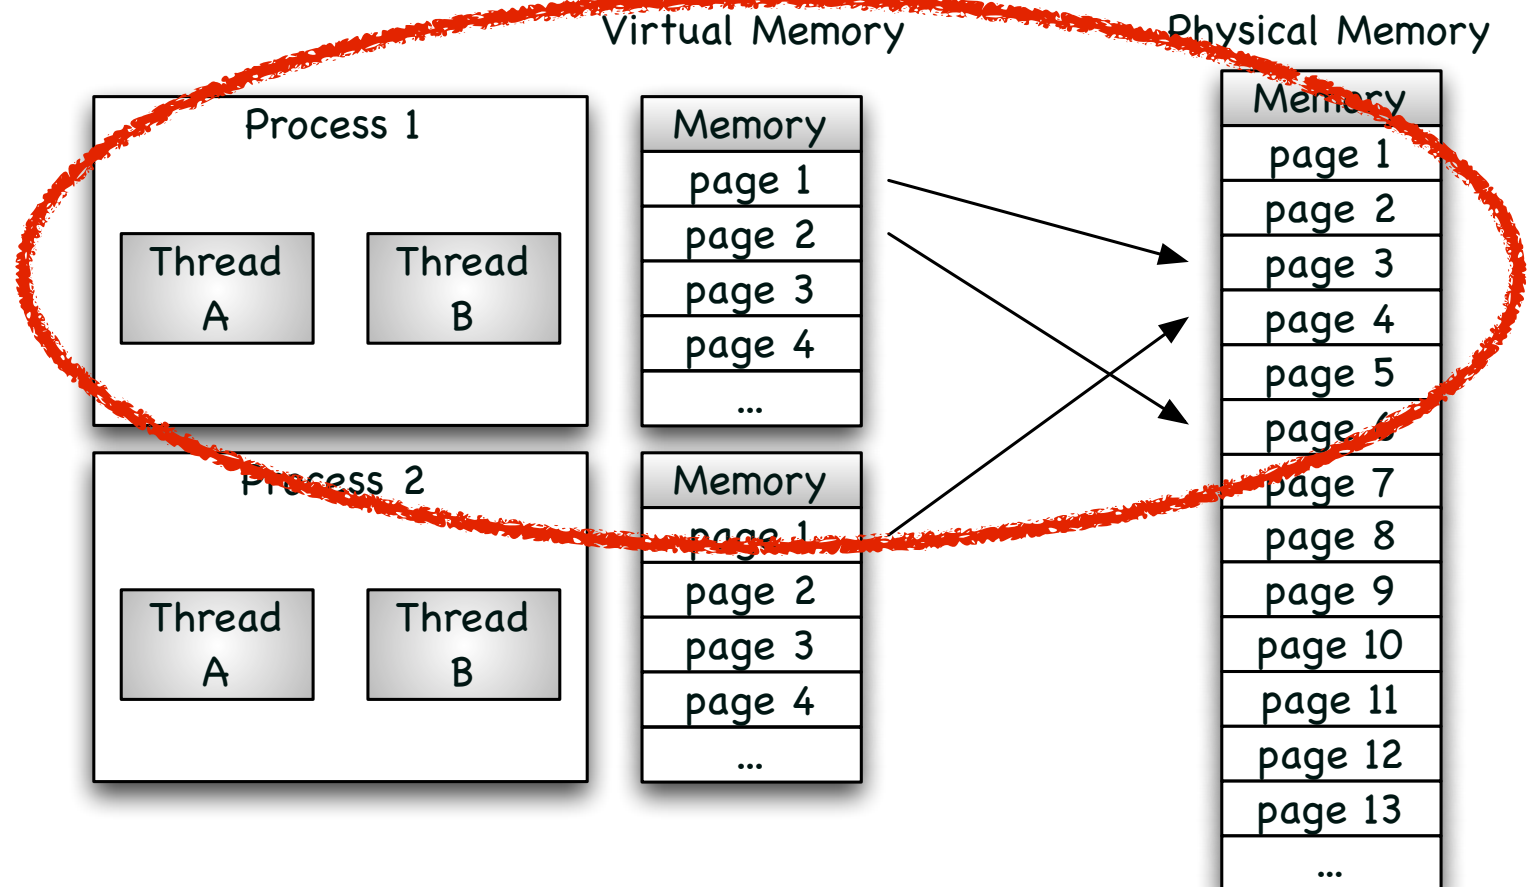
Parallel programs - A Solution

- *Parallel programs* are programs that are divided into execution units that can execute concurrently
 - ▶ Processes
 - ▶ Threads, jobs or tasks

Process & Threads in Linux

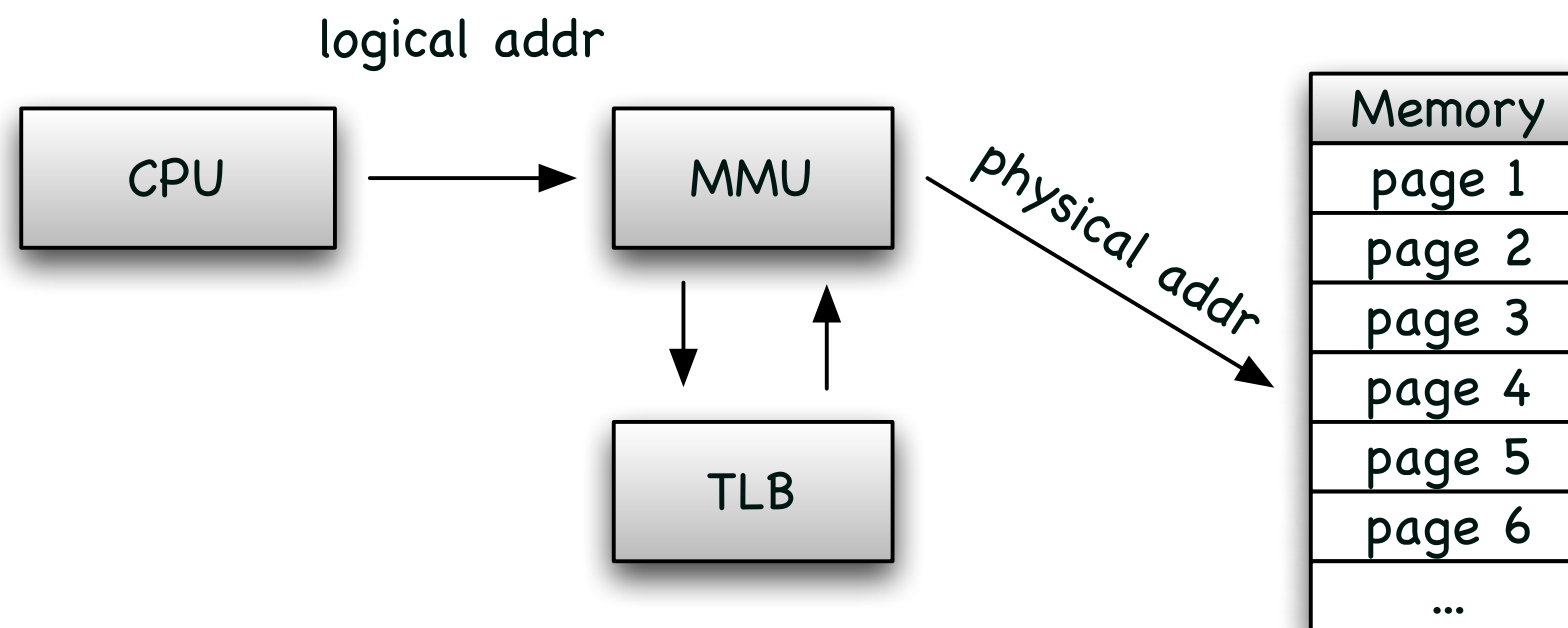
Processes and the OS (Kernel)

- Threads, Processes and Memory mapping
 - ▶ Each process has its own memory space
 - ▶ A mapping exists between virtual and physical memory
 - ▶ Not possible for one process to write in another address space
 - ▶ Threads share data space
 - ▶ Care must be taken *NOT* to destroy other threads data



Processes and the OS (Kernel)

- MMU - Memory Management Unit

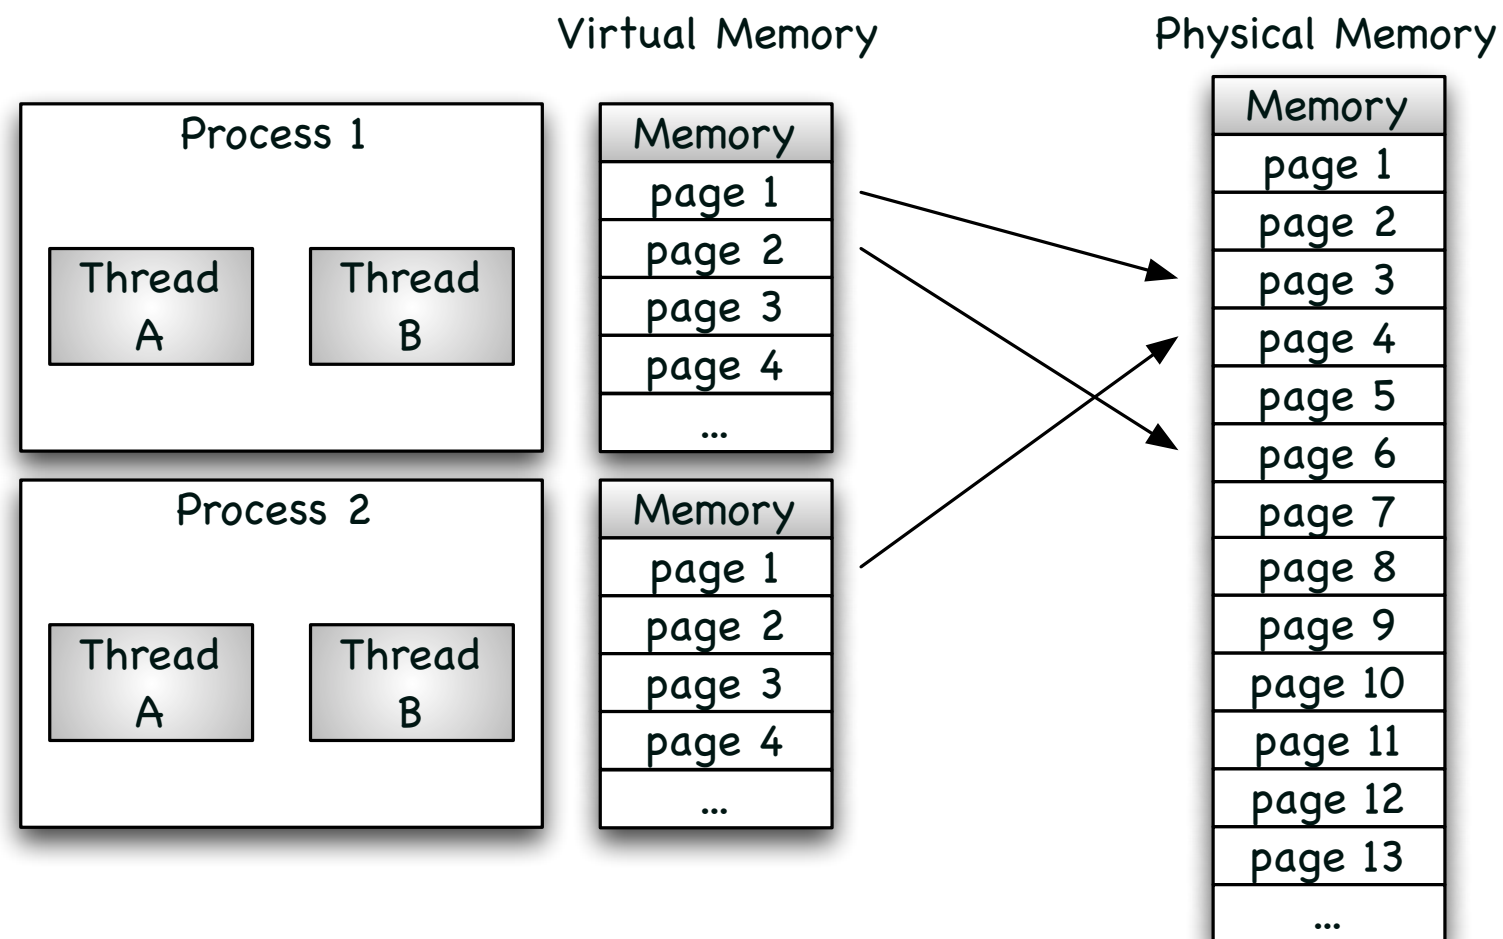


MMU: Memory Management Unit
TLB: Translation Lookaside Buffer

More to follow in MPS later

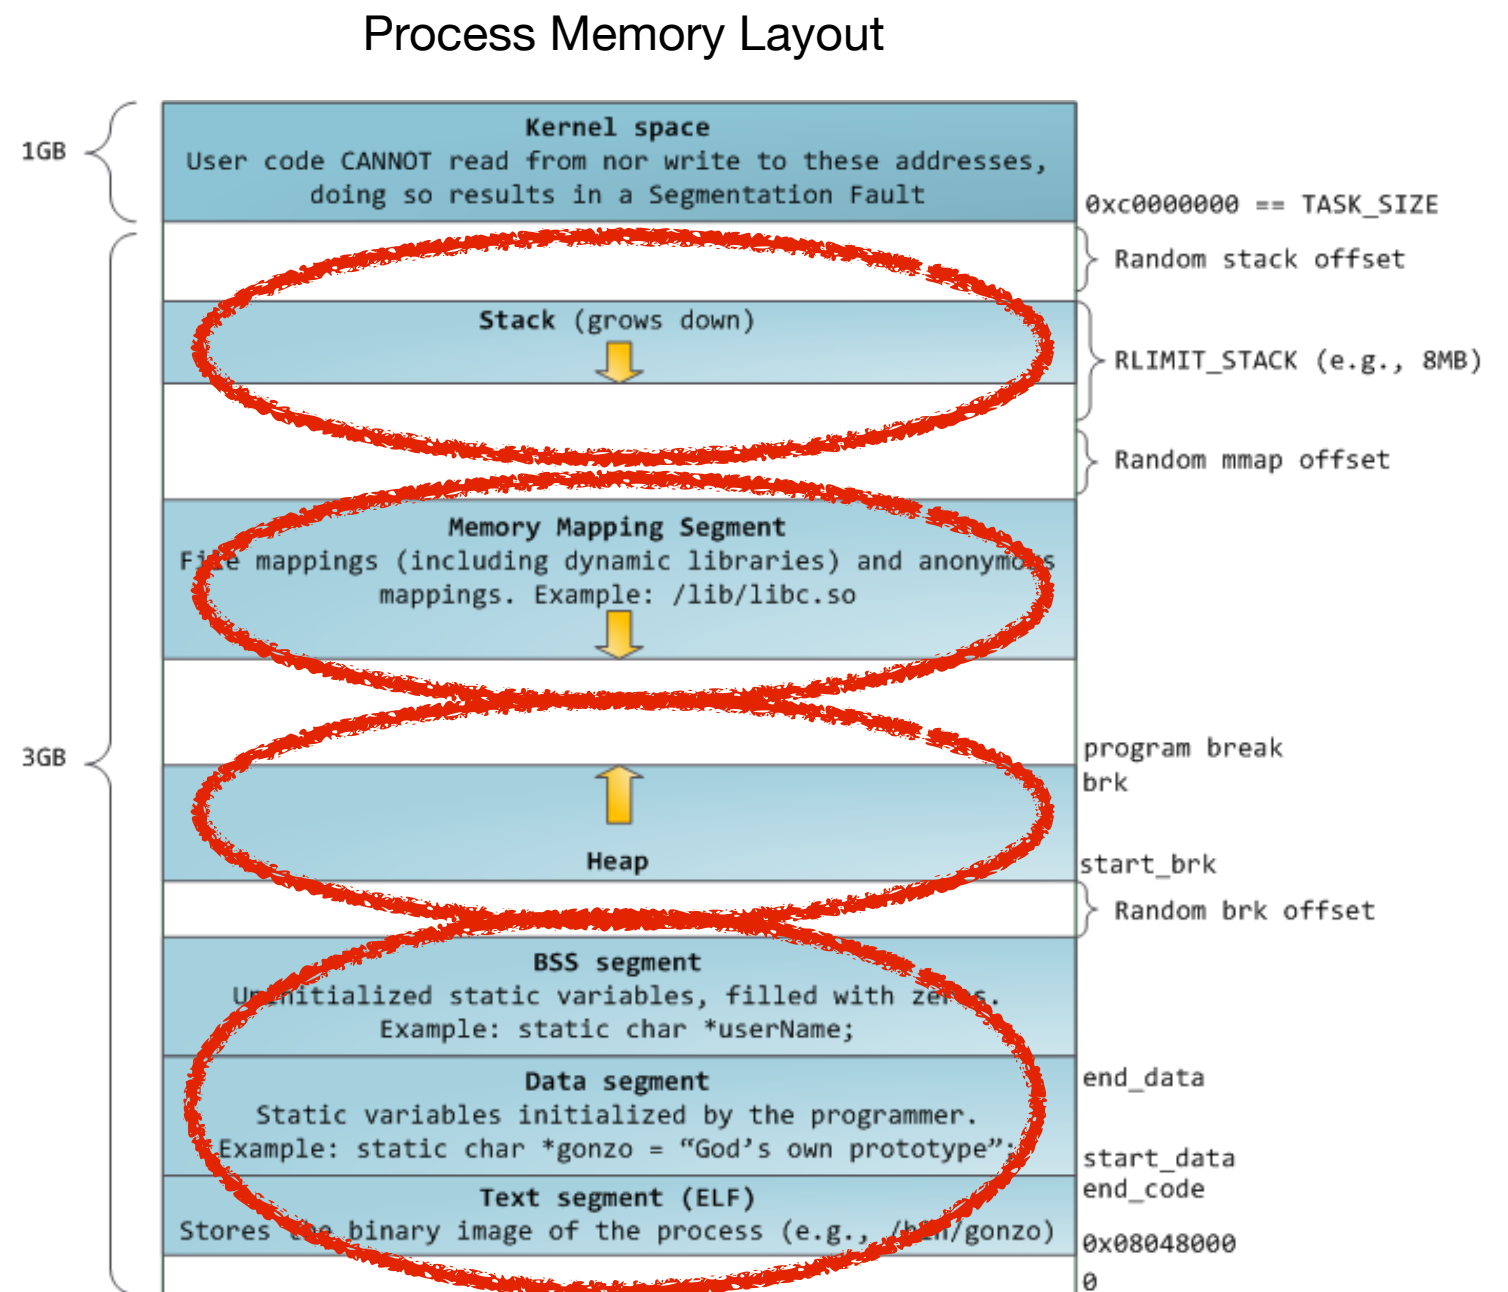
Processes and the OS (Kernel)

- MMU - Memory Management Unit



Process - What is it?

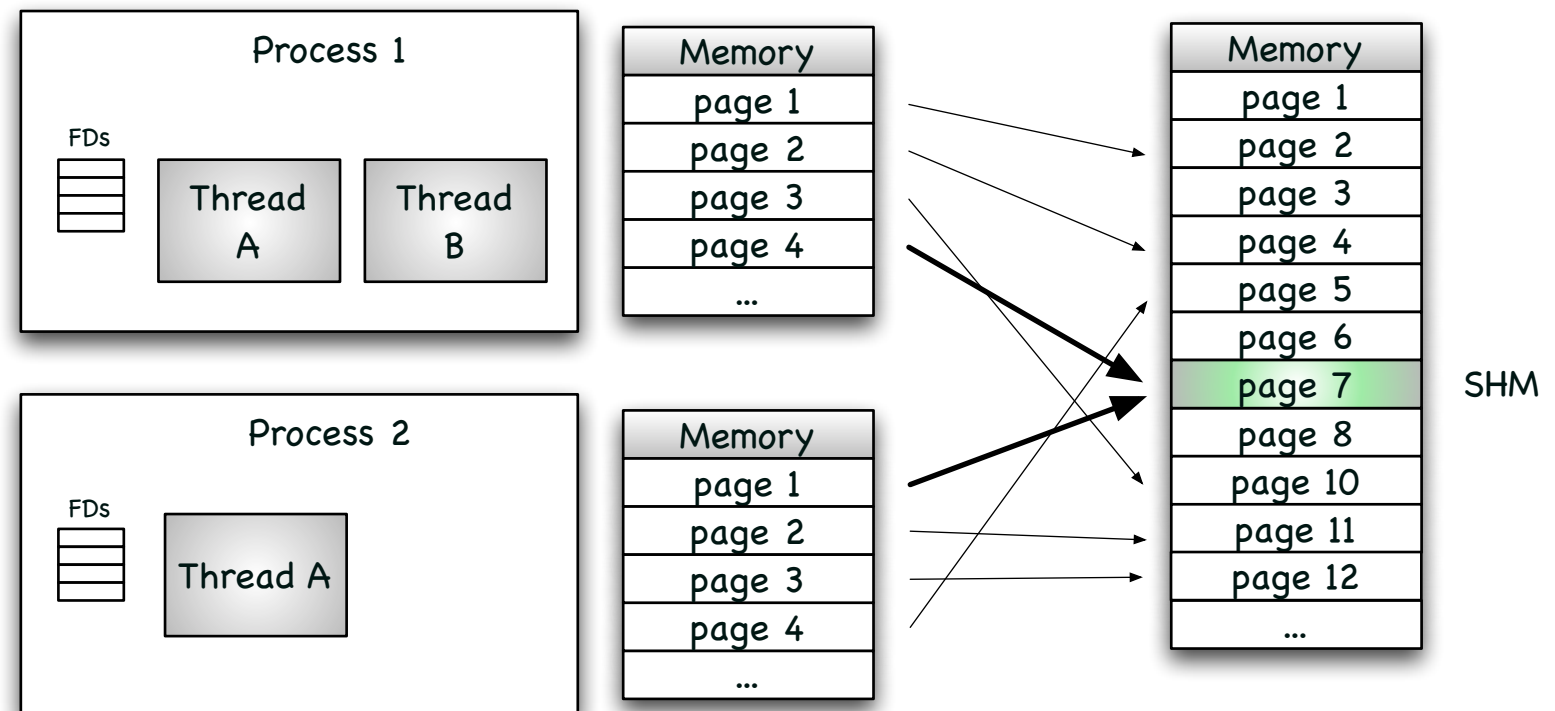
- Process - A program being executed in Linux
 - ▶ Stack
 - ▶ Local variables
 - ▶ Function return values
 - ▶ LIFO
 - ▶ Heap
 - ▶ “Free-store”
 - ▶ Dynamically allocated memory
 - ▶ Memory Mapping
 - ▶ File mapped in memory
 - ▶ Includes dyn libs
 - ▶ Sharing memory between processes
 - ▶ Variables and ELF



<http://duartes.org/gustavo/blog/post/anatomy-of-a-program-in-memory>

Processes and the OS - Sharing memory - IPC

- POSIX Shared Memory
 - ▶ Accessing memory affects other process
 - ▶ Pro
 - ▶ Speed/Performance
 - ▶ Cons
 - ▶ Fragile
 - ▶ Death of process
 - ▶ Data must abide certain principles
 - ▶ Challenge ensuring synchronicity



Threads of execution and the OS (Kernel)

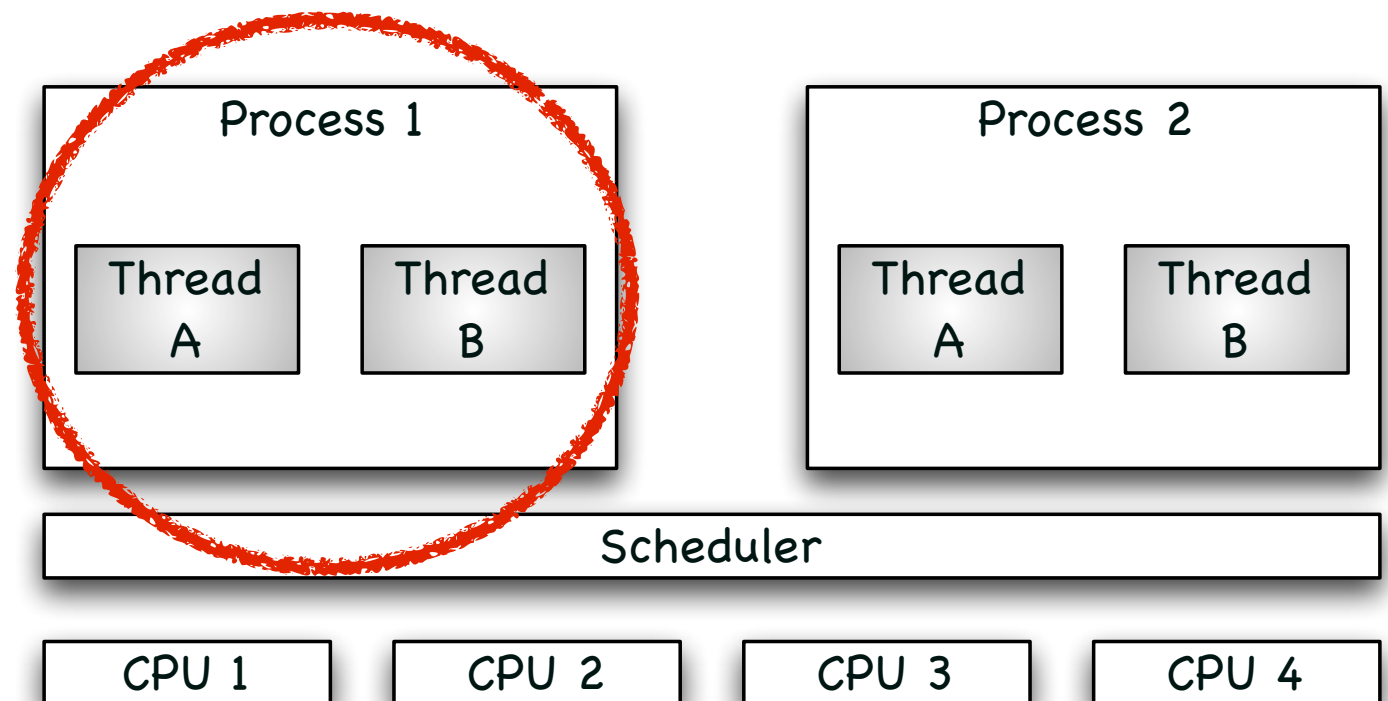
- How are threads mapped for execution?

- ▶ Three different models

- ▶ User level threading

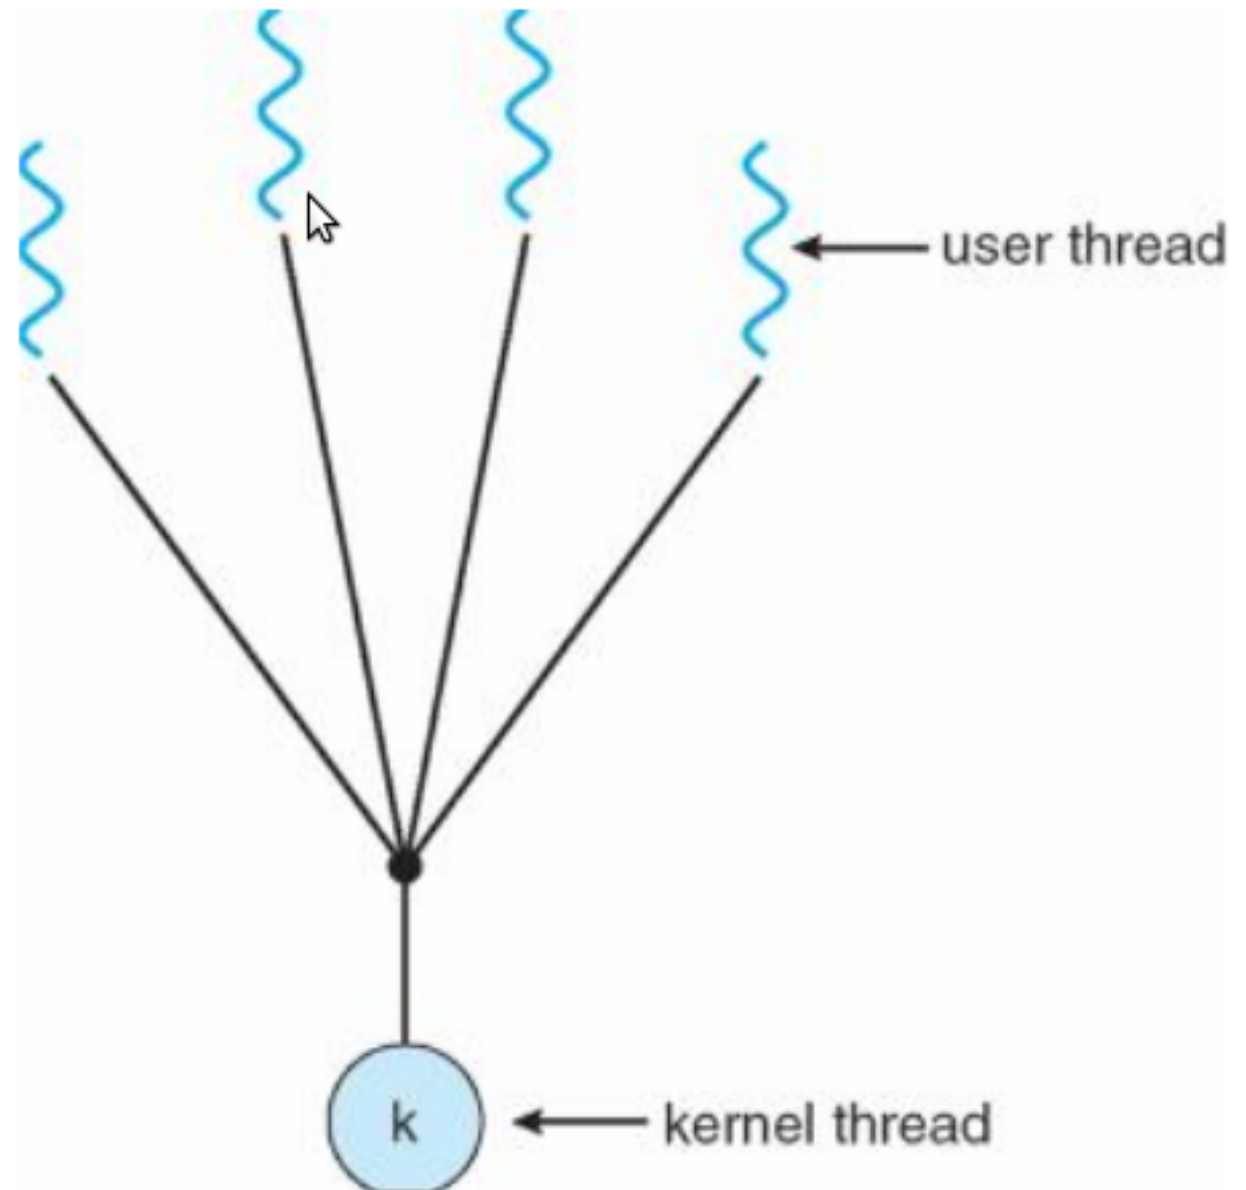
- ▶ Kernel level threading

- ▶ Hybrid level threading



Threading Model

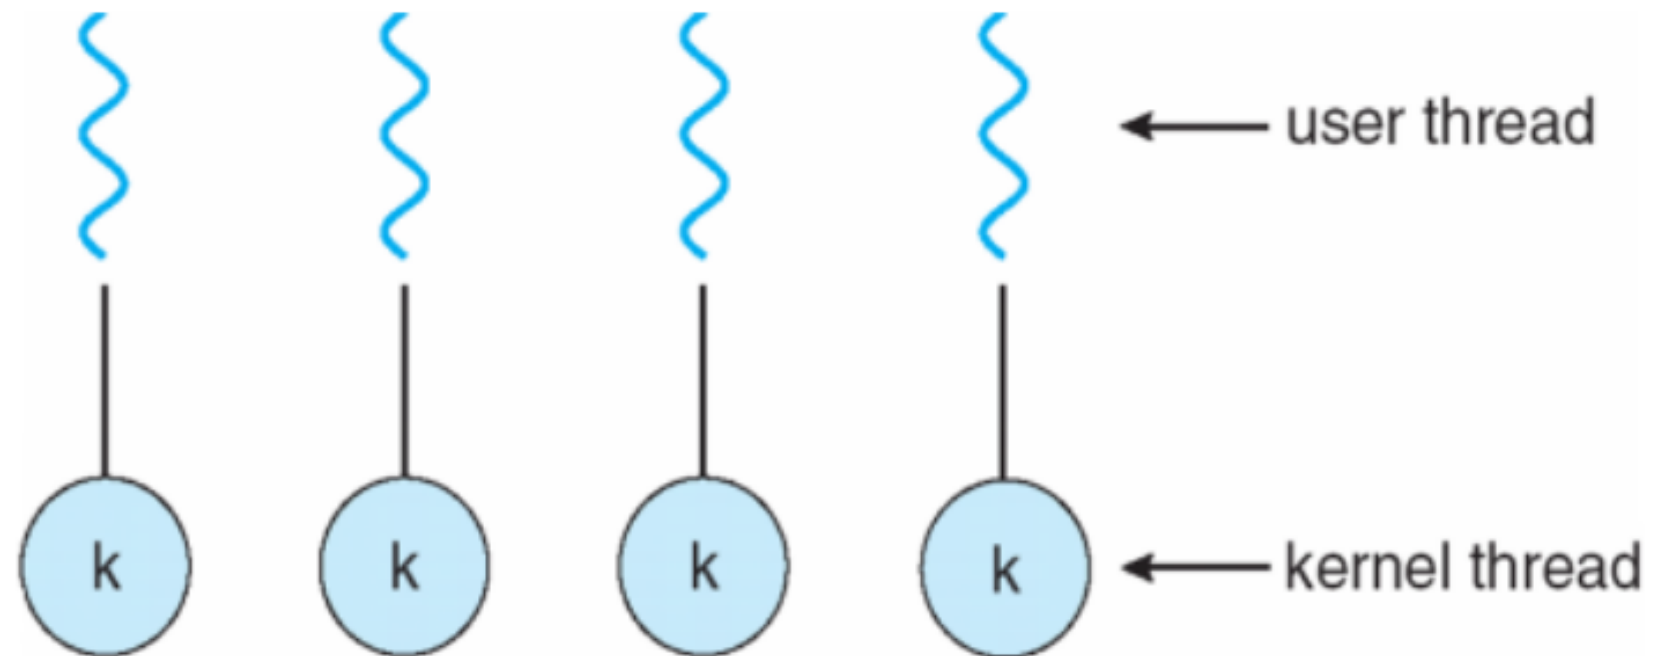
- User-Level Threading
 - ▶ Simple implementation no kernel support for threads
 - ▶ Very quick thread context switch (no kernel handling needed)
 - ▶ Not possible to handle multicores



Threading Model

- Kernel Level Threading

- ▶ Need thread awareness in kernel
- ▶ Maps directly to threads which the scheduler can control
- ▶ Efficient multicore usage

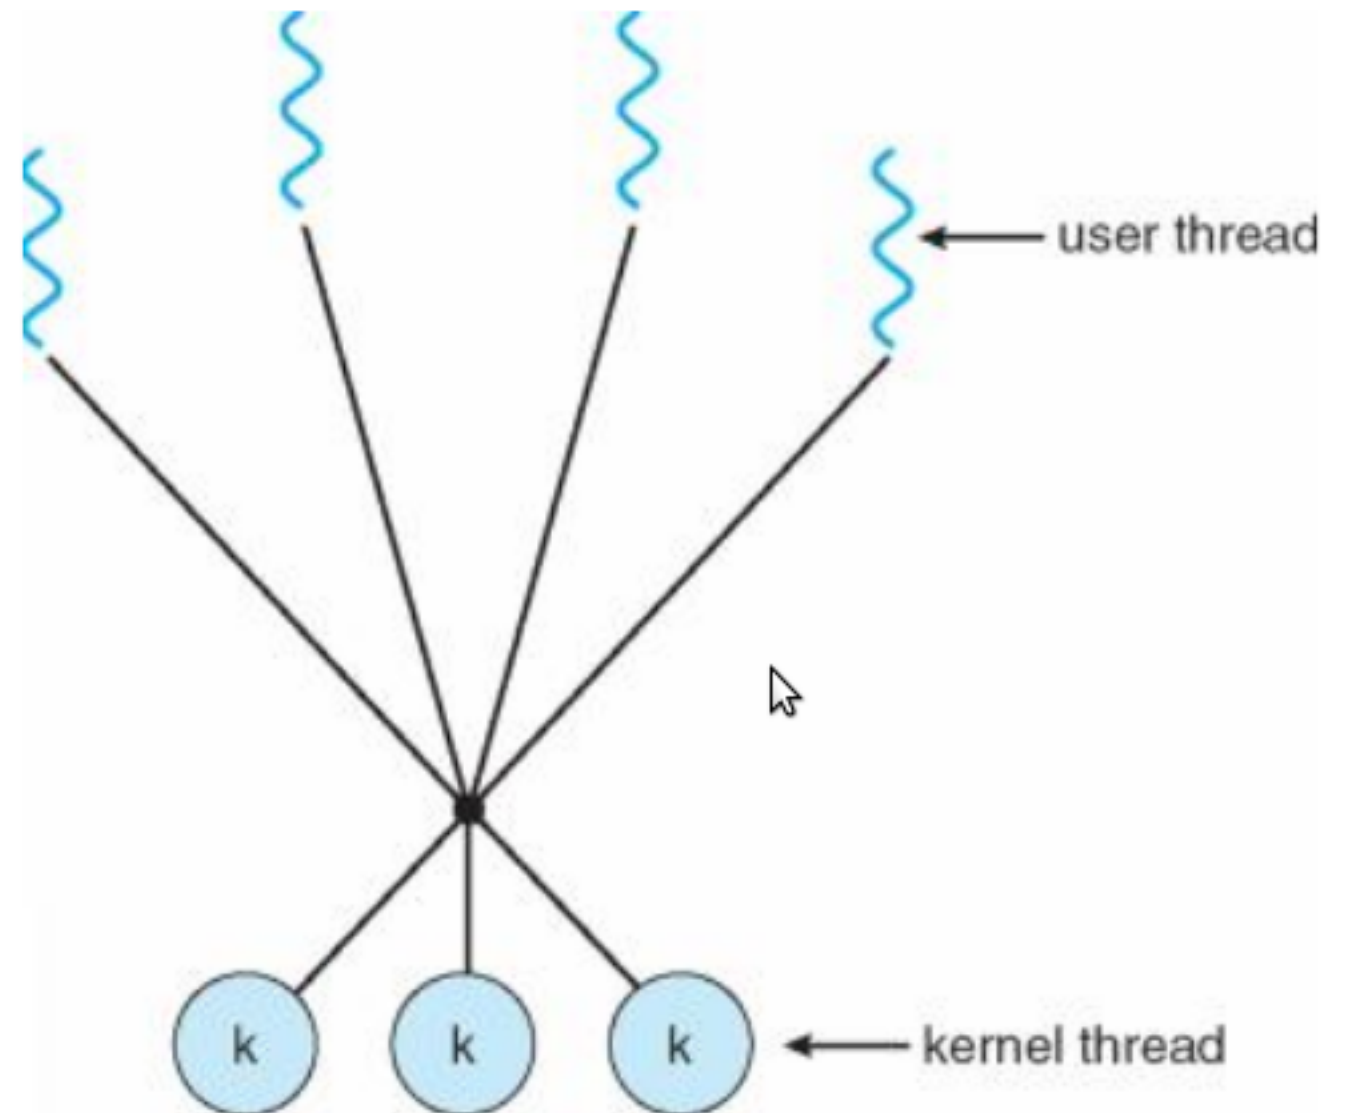


- OS

- ▶ Linux (NPLT), Win32 etc.

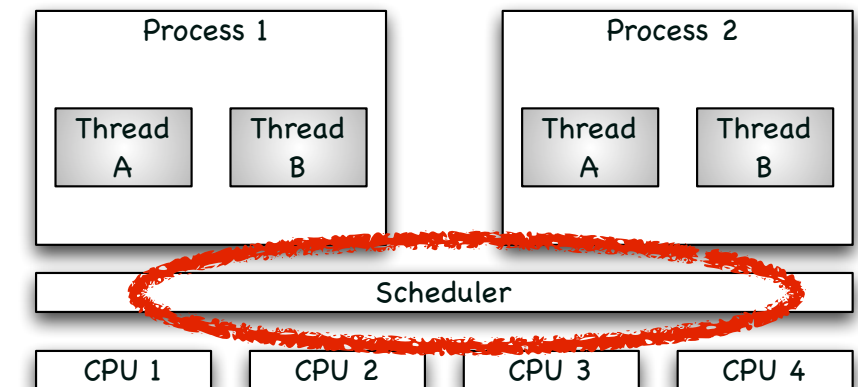
Threading Model

- Hybrid Level Threading
 - Complex implementation
 - Requires good coordination between user land and kernel land scheduler
 - Otherwise suboptimal resource usage
- OS
 - Windows 7



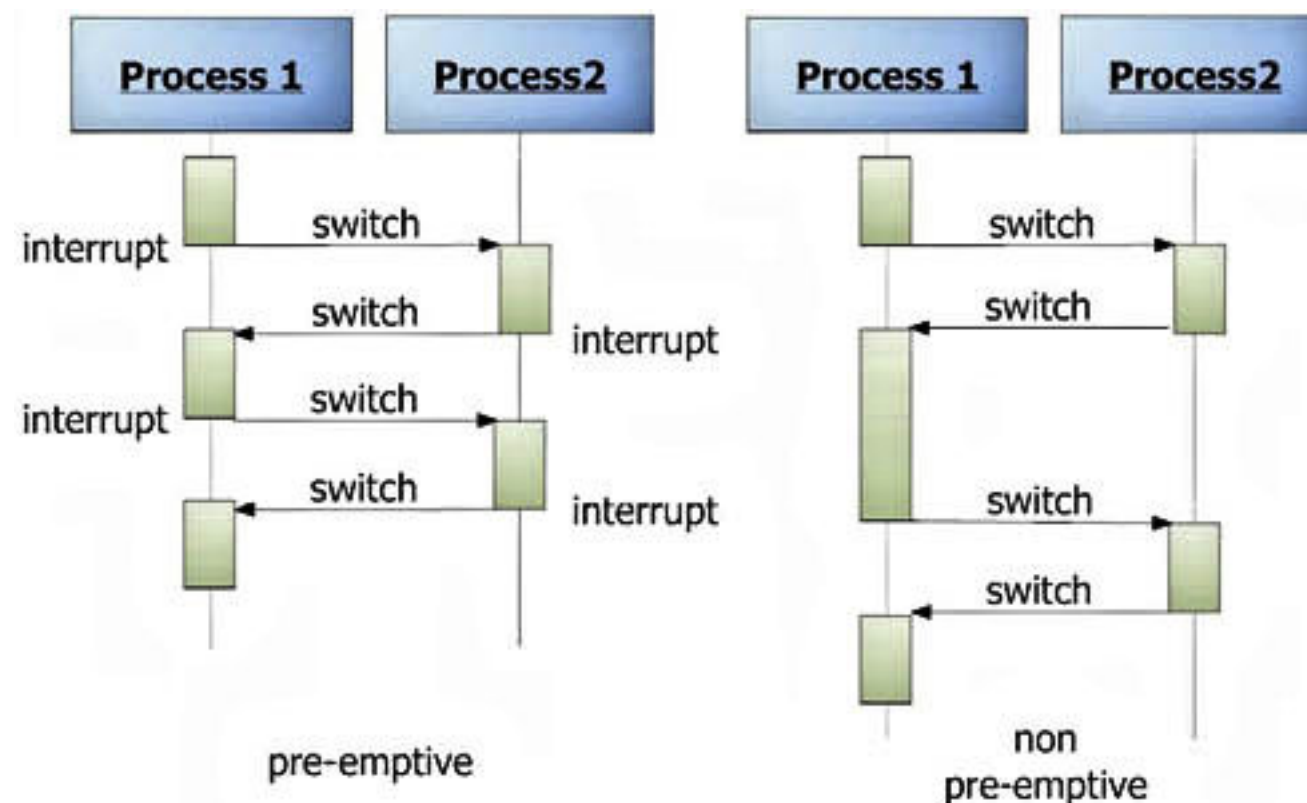
Context switching

- A *context* is the environment of the currently running process
- A context switch is performed by the OS to suspend the currently running process and resume another process
- General steps:
 - ▶ Interrupt current process
 - ▶ Save context of current process (SP, PC, registers, ...)
 - ▶ Restore context of next process
 - ▶ Resume execution of next process

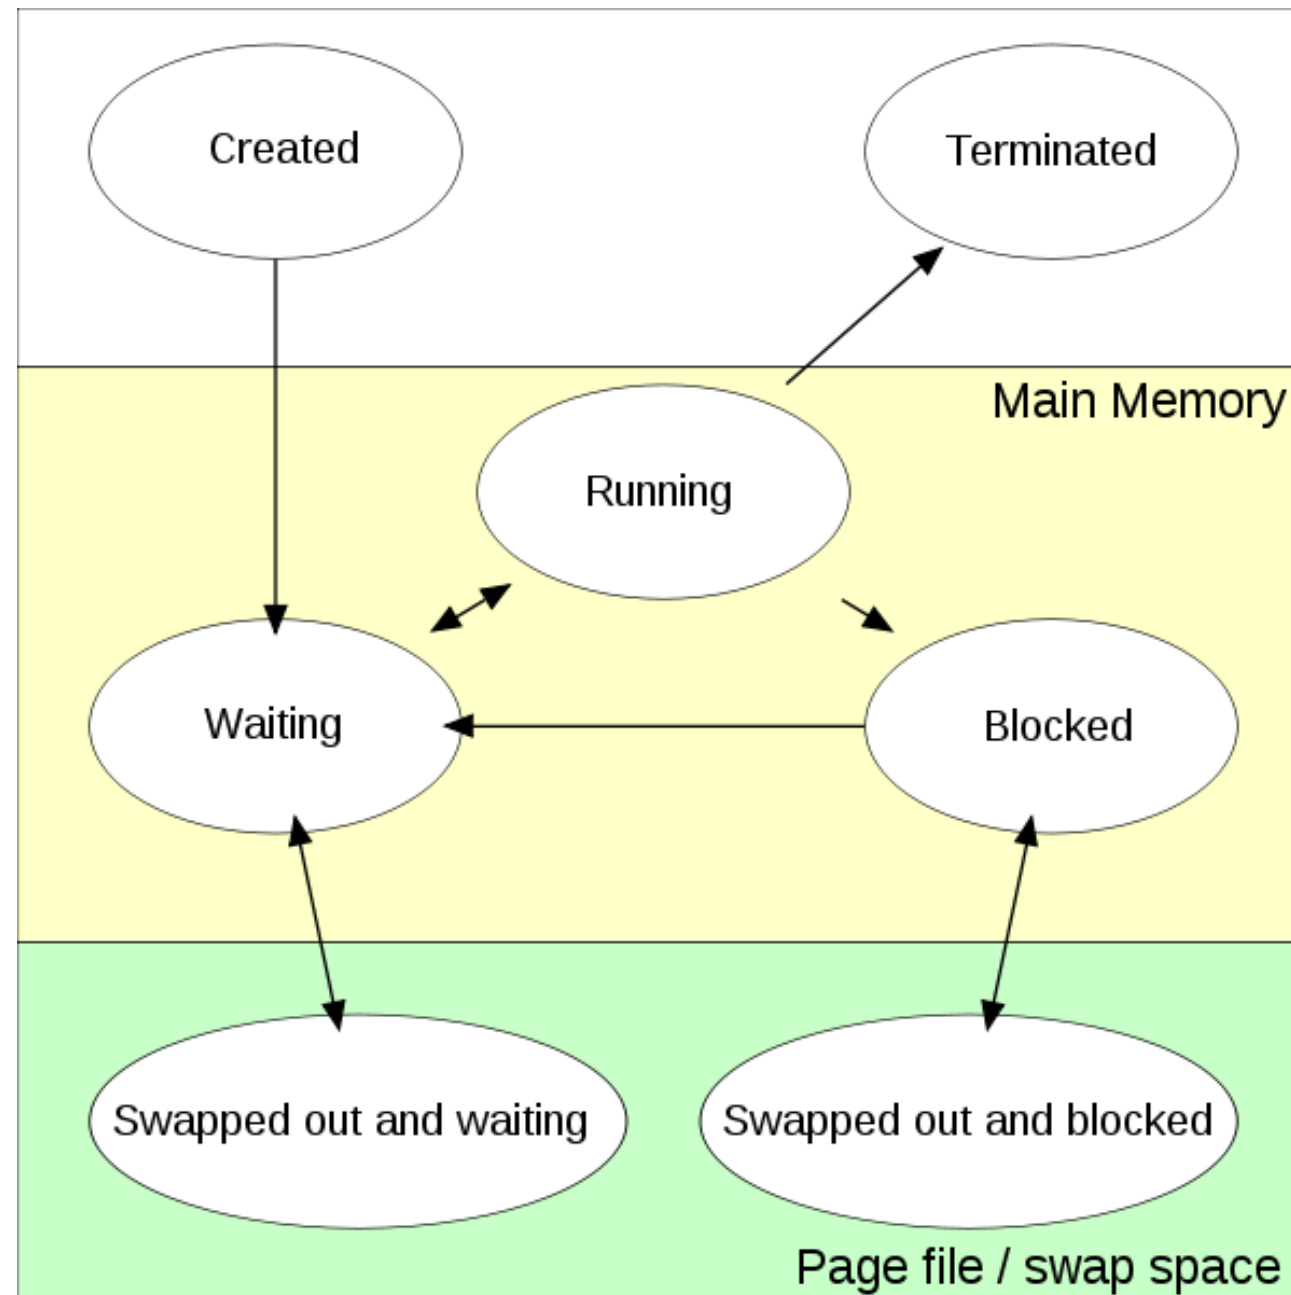


Context switching

- The *operating system* schedules tasks for execution
 - Preemptive scheduling: Tasks can be interrupted at any time
 - Nonpreemptive scheduling: Tasks voluntarily yield the CPU
- *Linux supports both - Kernel configuration options*



The life and death of a process



Threads - Summary

- Threads are *strands of execution* – each process has at least one thread
 - ▶ AKA *light-weight processes*
 - ▶ Also called tasks or jobs
- Threads of the same process share memory space (e.g. global variables)
- Threads *can harm* each other
- We will often work with several *threads* in the same *process*

Processes - Summary

- A process is an instance of a program that is being executed
 - ▶ Image of program (segment),
 - ▶ Stack, heap, registers, file descriptors, ...
- Processes have their own individual memory spaces
 - ▶ Process A cannot write in memory of process B – they are safe from each other
- Processes may only communicate through IPC mechanisms controlled by the OS.
- Processes may spawn other processes, which may execute the same or other programs
- A process may also spawn threads within its own memory space

Multitasking systems: Advantages

- What are the advantages of multiple tasking a system?
 - ▶ Prioritization – the highest-priority task gets to run
 - ▶ Modularization – wrap concurrent activities in a task
 - ▶ Resource utilization – Don't spend CPU time waiting for I/O etc.
 - ▶ ...
- However, the use of multiple tasks in a system creates a number of problems for us
- We must know the problems to be able to identify and counter or avoid them

Multithreaded systems: Disadvantages

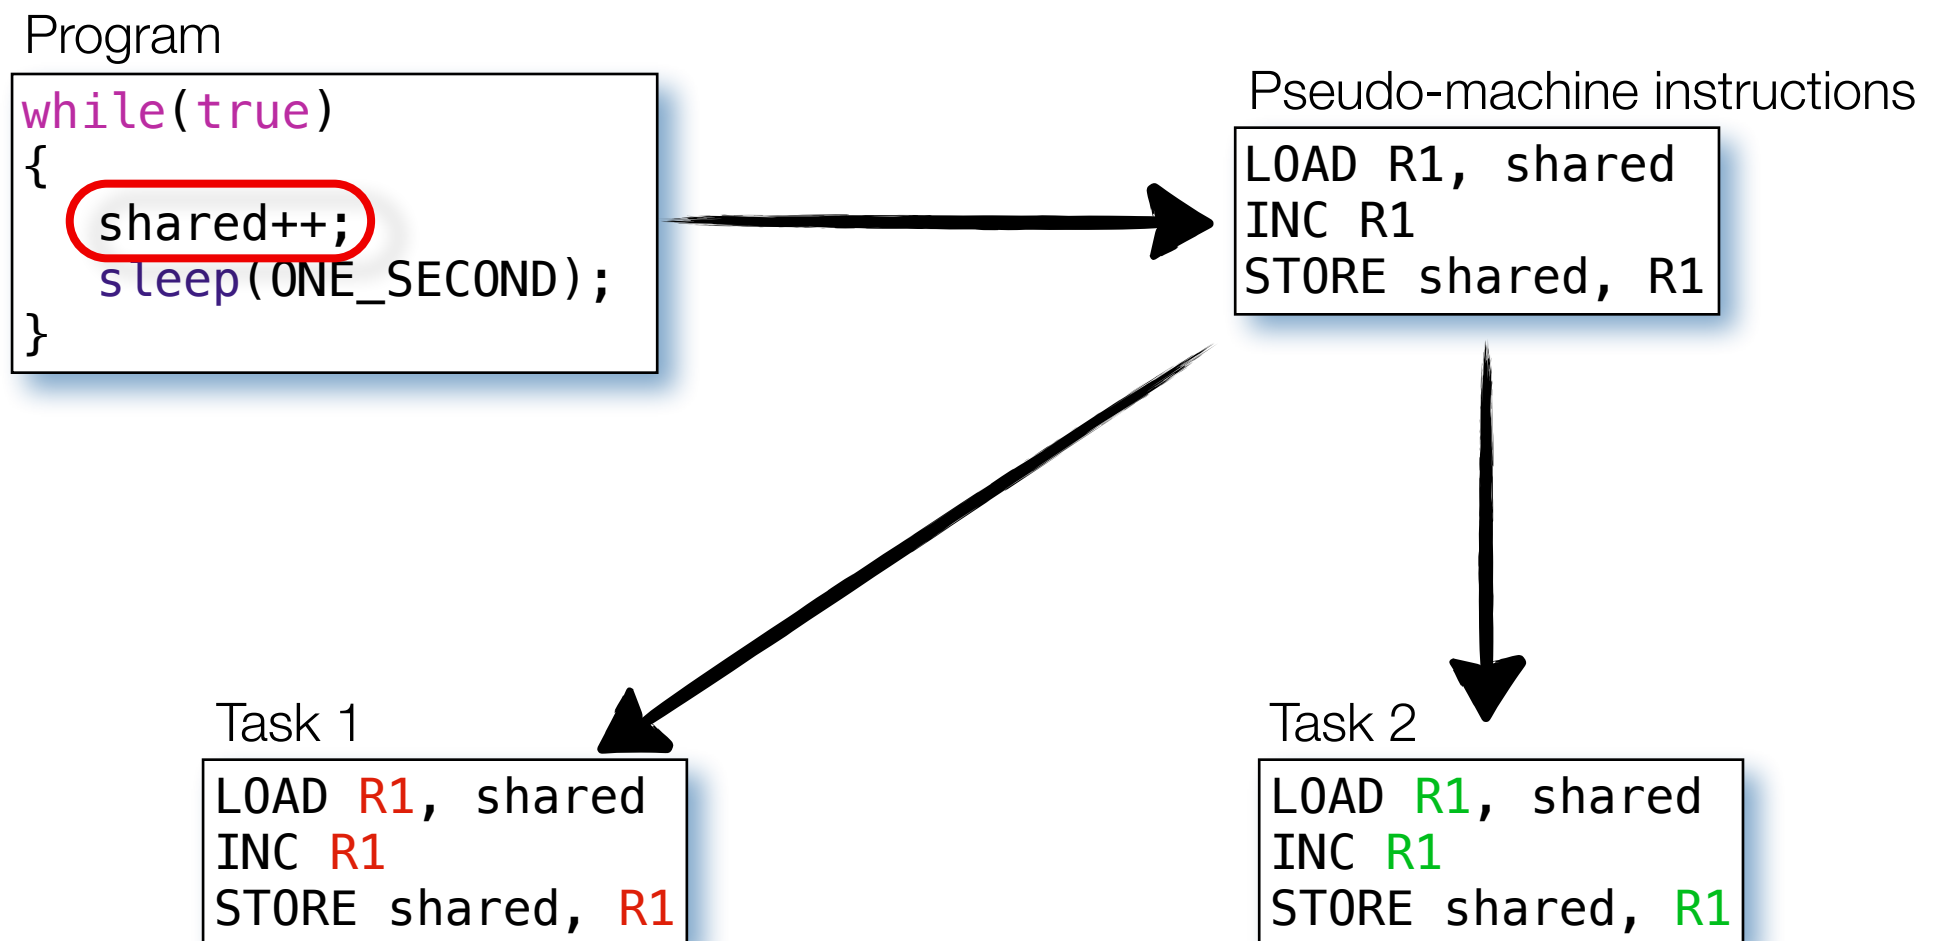
- Consider the following code. What is the value of shared after 10 seconds?

```
unsigned int shared;
void taskfunc()
{
    while(true)
    {
        shared++;           // Increment i, then wait
        sleep(ONE_SECOND);  // 1 second
    }
}

int main()
{
    shared = 0;
    createThread(taskFunc); // Start two identical threads
    createThread(taskFunc); // that run the same function
    while(true)
        sleep();
}
```

The shared data problem

- Let's zoom in:



The shared data problem

Task 1

```
LOAD R1, shared
INC R1
STORE shared, R1
```

Task 2

```
LOAD R1, shared
INC R1
STORE shared, R1
```

What value is shared
suppose to have?

Non-interleaved instructions

```
LOAD R1, shared // shared = 0
INC R1
STORE shared, R1 // shared = 1
LOAD R1, shared // shared = 1
INC R1
STORE shared, R1 // shared = 2
```

Which values can shared
in fact have?

Interleaved instructions

```
LOAD R1, shared // shared = 0
LOAD R1, shared // shared = 0
INC R1
STORE shared, R1 // shared = 1
INC R1
STORE shared, R1 // shared = 1
```

The shared data problem

```
Position pos { 10, 20, 30 };
```

Task 1

```
void newPos(Position& ps, float x, float y, float z)
{
    pos.x = x;
    pos.y = y;
    pos.z = z;
}
```

Task 2

```
void printPos(Position& pos)
{
    std::cout << "X: " << pos.x << std::endl;
    std::cout << "Y: " << pos.y << std::endl;
    std::cout << "Z: " << pos.z << std::endl;
}
```

Non-interleaved instructions

```
T1 pos.x = x;
T1 pos.y = y;
T1 pos.z = z;

T2 std::cout << "X: " << pos.x << std::endl;
T2 std::cout << "Y: " << pos.y << std::endl;
T2 std::cout << "Z: " << pos.z << std::endl;
```

newPos called with 11, 22, 33

std::cout prints
X: 10
Y: 20
Z: 30

Interleaved instructions

```
T1 pos.x = x;

T2 std::cout << "X: " << pos.x << std::endl;
T2 std::cout << "Y: " << pos.y << std::endl;
T2 std::cout << "Z: " << pos.z << std::endl;

T1 pos.y = y;
T1 pos.z = z;
```

Pseudo'ish code

std::cout prints
X: 11
Y: 20
Z: 30

Hey WHAT???

The shared data problem


- The shared data problem is inherent in any preemptive multithreaded system
- Very cumbersome to find a good software solution to the problem
 - Peterson's solution: 2 interest flags, 1 will-wait flag
 - Does not scale
- We need a way to define critical sections of program
 - Sections in which the thread is guaranteed to be allowed to execute uninterrupted
- In a later lecture!

```
void taskfunc()
{
    while(true)
    {
        enterCriticalSection();
        shared++;
        exitCriticalSection();
        sleep(ONE_SECOND);
    }
}
```


Multithreaded systems: Disadvantages

- Consider a system with three threads, HP, MP and LP:

- ▶ HP takes 5 μs , must run once every 10 μs
- ▶ MP takes 3 μs , must run once every 10 μs
- ▶ LP takes 1 μs , must run once every 1000 μs




- Schedule?




OK!

Starvation

- Assume we add another MP thread, MP_2
 - MP_2 takes 2 μs , may run once every 10 μs



- Schedule?



OK!

Starvation!

Starvation

- Starvation is an inherent problem in any priority-based system
- It occurs when the schedule is so tight that LP threads are never allowed to run because higher-priority threads “hog” the CPU
- Starvation can be very hard to predict and detect
 - ▶ Might only occur in very special situations

Programming with threads

- C++ does not have the concept of threads built in
 - ▶ Relies on 3rd party libraries for this
 - ▶ *Not entirely true see new ratified C++ standard*
- The POSIX (Portable OS Interface for uniX) library is the most widely used threading library
 - ▶ Others include boost, Qt, ...
- The POSIX library has the thread type pthread which we will use.
 - ▶ Include pthread.h, link with library pthread

PThread functions

- Family of pthread functions you are to use...
 - ▶ pthread_create()
 - ▶ pthread_join()
 - ▶ pthread_exit()
 - ▶ pthread_* (and more)
- How they work is for next session