

Embedded Software

Inter-thread communication (Intra-process communication)

Agenda

- Communication design challenges
- Message Queue and Handler Design
- Consequences

Communication design challenges

- Individual threads wait for a condition to become true
- Enter and leave critical sections using mutexes or semaphores
 - May happen multiple times in the space of one thread loop iteration
- May even hold multiple resources which have to be synchronized between threads
 - The sequence in which resources are taken must be thought through.
- **Consequence**
 - A design challenge ensuring that no deadlocks or timing issues exist
 - Readability easily becomes an issue too
 - High code complexity is the outcome

What we need...

- We want an approach where
 - ▶ *all* processing within a thread must *not* require locking
 - ▶ however *other* threads must be able to pass control and/or data to a specific thread via some mechanism.
 - ▶ *multiple* threads may concurrently decide to pass such control and/or data

A step backwards

- What is it in fact we are doing and what?
 - ▶ *Perform some action when a given condition becomes true or we get signaled*

We want events (messages)!

Event Driven Programming


- Reactionary programming
 - ▶ ***Each incoming message is processed by a specific handler***
 - ▶ E.g. its a handler - *it reacts someone must take initiative!*
 - ▶ Types
 - ▶ Sensor input
 - ▶ Temperature exceeded message → Turn down heat
 - ▶ Car detected wanting to enter car park message → Open garage door
 - ▶ Signal input
 - ▶ Exit button in GUI message → Exit program

Event Driven Programming (Event = Message)


- Can be viewed as a two phase process
 - ▶ Acquire/Select new message
 - ▶ Handled by a *Message Queue* and ensures that a number messages can be in “queue” at a time
 - ▶ Process new message in handler
 - ▶ Handled by casing out on the specific message

Message Queue & Handler design

Resembles the “Producer & Consumer problem”



- The producer-consumer problem
 - ▶ A producer thread produces buffer items
 - ▶ A consumer thread consumes them
- Applied to our problem we get



Further requirements for *our* Message Queue


- If the receiving queue is full, then the thread or threads wishing to pass control and/or data must block waiting for more space.
 - ▶ Implies that there *is* a maximum number of elements in a queue
- The consuming thread *must block* upon receiving from an empty queue
- Blocks are NOT to be done with polling (+ sleeps), *why?*
- What should we do then? - ***Conditionals***

What is the structure of the information to pass around?

- void* or simple array of bytes
 - ▶ Can contain anything
 - ▶ No type information - No type-safety (if we don't know what it is - we don't know how to delete)
- template based
 - ▶ Depends on the implementation, is a good solution but more complex
 - ▶ Type-safety
- Inheritance
 - ▶ Simple and extended via sub-classing
 - ▶ Type-safety / Type information - Delete via base pointer
 - ▶ Might incur overhead

Inheritance - our choice

- Message hierarchy




```
class Message
{
public:
    virtual ~Message(){}
};
```

```
struct GarageDoorOpenReq : public
Message
{
    MsgQueue* mq_;
};
```

Message Parsing

- A producer creates and “sends” a GarageDoorOpenReq message
 - ▶ class GarageDoorOpenReq is therefore seen as a message

How does the receiver determine which message it is?



From parent to child

- How do we convert a *Message** to a *GarageOpenDoorReq**?

- ▶ Via using `dynamic_cast<>`

```
GarageDoorOpenReq gdor;  
Message* msg_ = &gdor; // Illustration!
```

```
GarageDoorOpenReq* req = dynamic_cast<GarageDoorOpenReq*>(msg_);  
// Runtime check, req == NULL if not correct
```

- ▶ Via `typeid()`


```
GarageDoorOpenReq gdor;  
Message* msg_ = &gdor; // Illustration!
```

```
if (typeid(*msg_) == typeid(GarageDoorOpenReq))  
{  
    // Runtime check – evaluates to true if pointer is of said type  
    GarageDoorOpenReq* req = static_cast<GarageDoorOpenReq*>(msg_);  
}
```

From parent to child

- How do we convert a *Message** to a *GarageOpenDoorReq**?
 - ▶ Using a special identifier
 - ▶ associating an id with the message

```
enum
{
    ID_GARAGE_DOOR_OPEN_REQ=0,
    ID_GARAGE_DOOR_OPEN_CFM=1,
    ID_XXX=2,
    ID_YYY=3
};
```



Considerations regarding Embedded Systems

Embedded Compiler configurations

- However certain embedded compilers are compiled without support for RTTI and exception.
 - ▶ RTTI - Run Time Type Information
 - ▶ Costs in the form of space - *Yes it costs, but what are the consequences?*
 - ▶ Exceptions
 - ▶ The perception is:
 - ▶ Costs in the form of space - *What would the code handling normal errors costs?*
 - ▶ It is difficult to do correctly - *Thats certainly correct, but it is not impossible*
 - ▶ Errors are not tolerated at all, they must all be found - ***That is*** If you have the time and money, depends on the amount money


Embedded Compiler configurations

- Based on these inputs the following requirement is added:
 - ▶ *It is acknowledged that the use of RTTI will improve program readability, however due to the increase in code size it is denounced*
 - ▶ Meaning no use of: (in our design)
 - ▶ `dynamic_cast<>` - Runtime check whether the cast is permissible or not
 - ▶ `typeid()` - Uniquely identify a given object

Due to compiler considerations

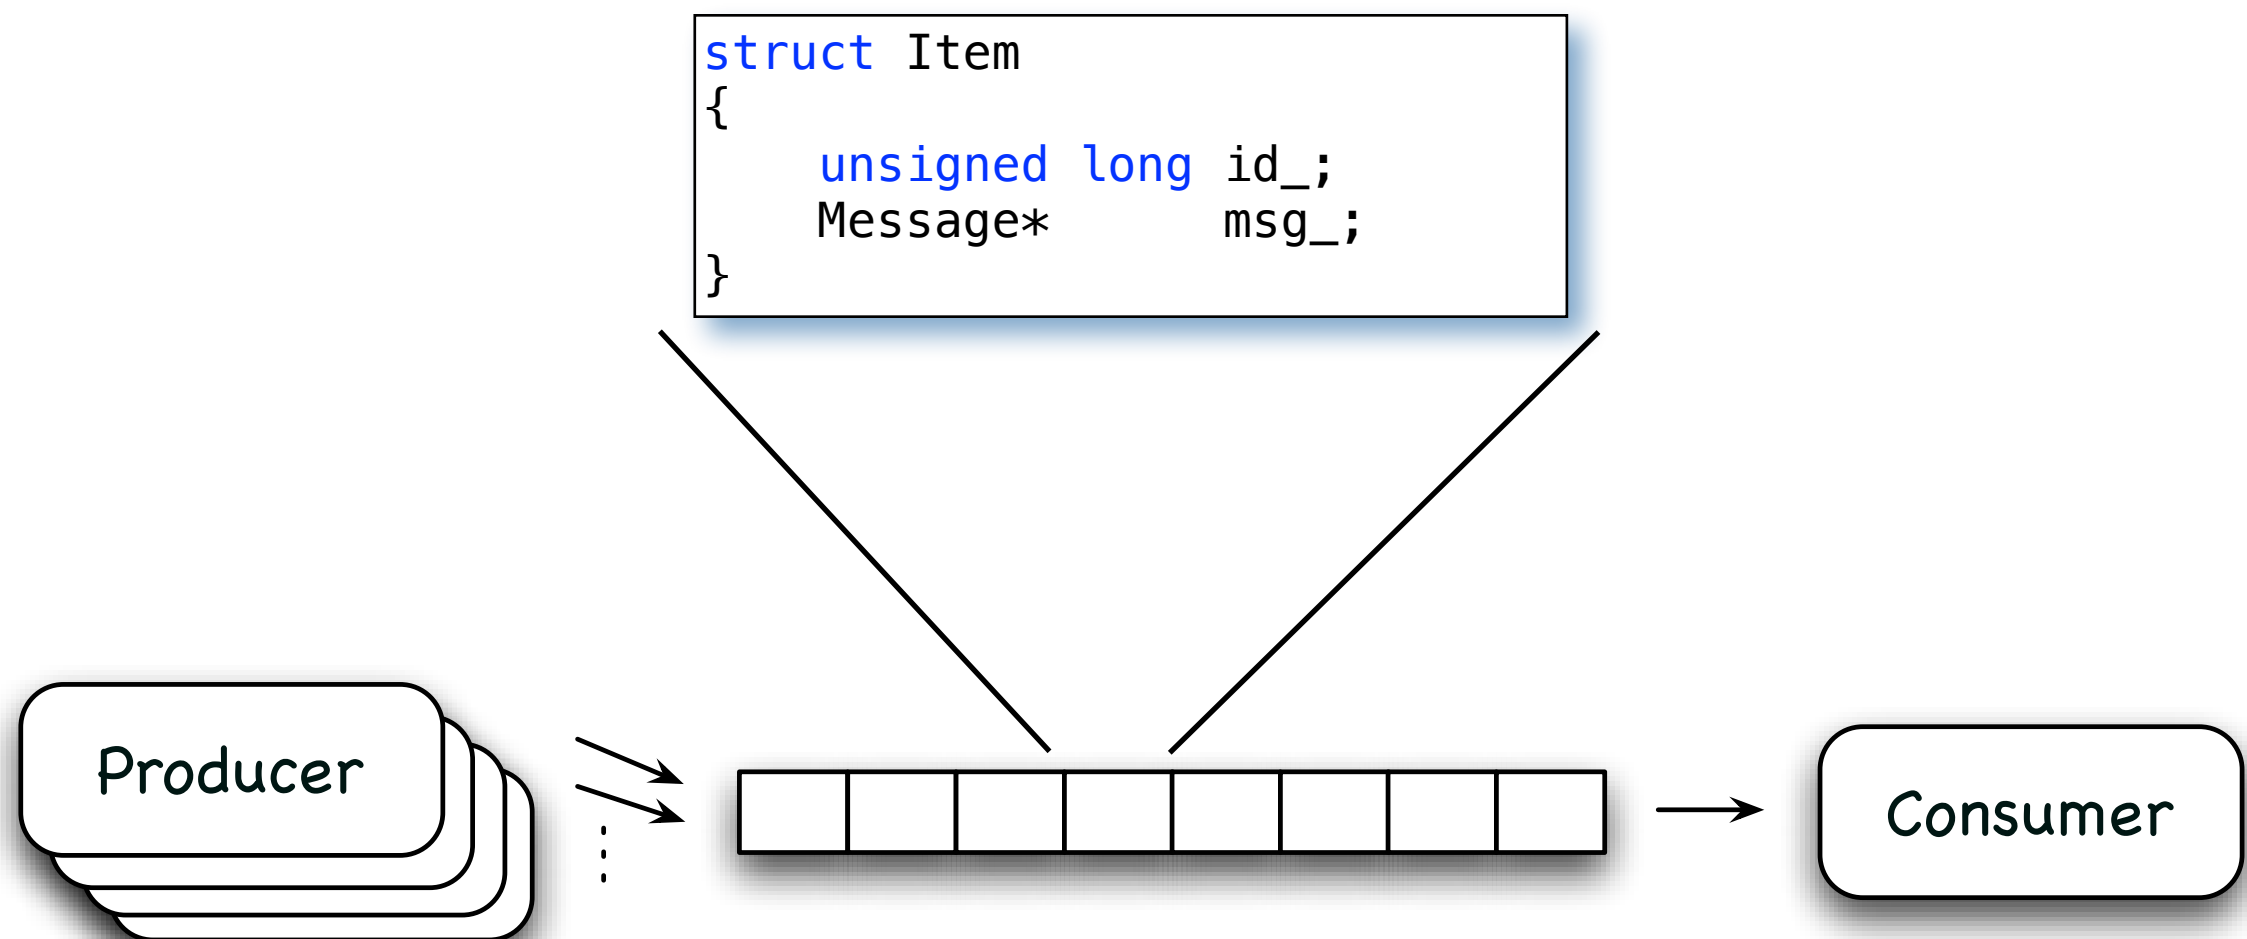
- We will be associating an id with the message

```
enum
{
    ID_GARAGE_DOOR_OPEN_REQ=0,
    ID_GARAGE_DOOR_OPEN_CFM=1,
    ID_XXX=2,
    ID_YYY=3
};
```

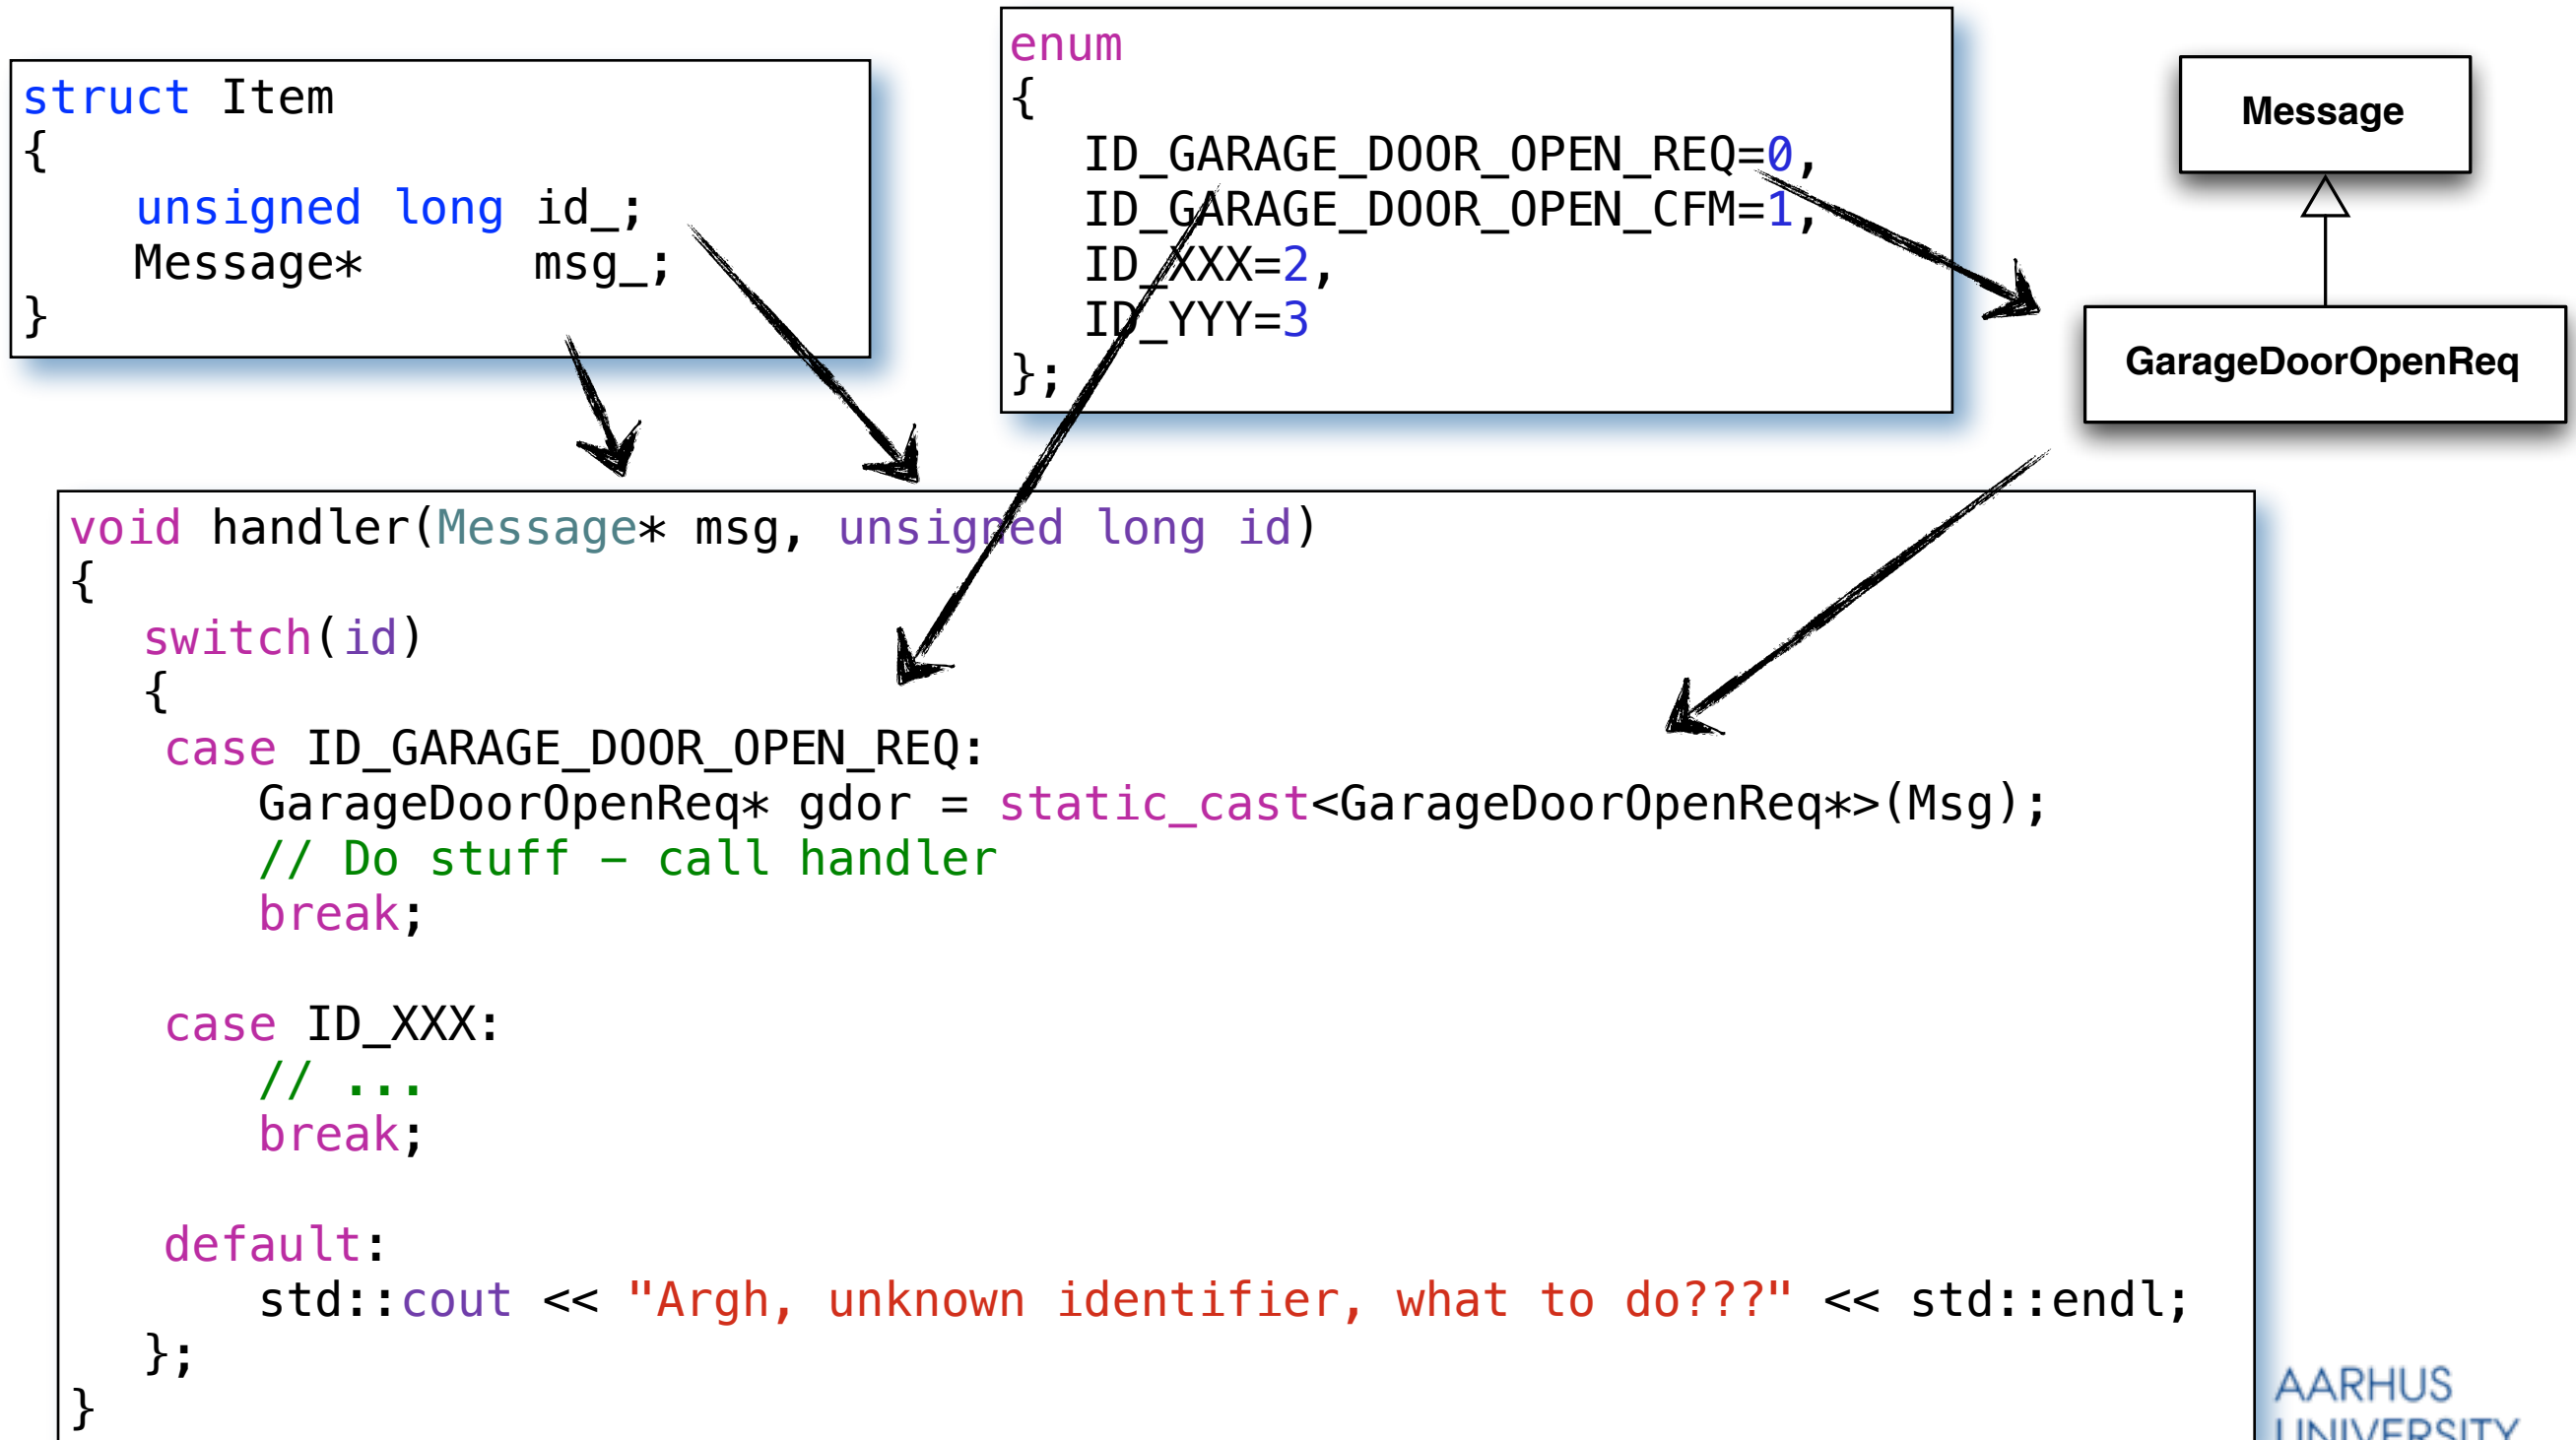


Choice of item in MsgQueue

- id_ is the identifier which is to be send
- msg_ is the message to be passed



An identifier to designate which child it is (the handler)



Message / ID combo

- Associate an identifier with a class/structure
 - ▶ The compound signifies the control/data information to be send/received
 - ▶ The identifier is denoted by the receiving party NOT part of a globally defined enum; ***why not? Placed in a central place everyone knows; seems very good...?!***

The desired MsgQueue interface design

Sender threads use ***send()*** function to send messages to thread

MsgQueue
- queue_ : std::xxx
- maxSize_ : unsigned long
+ MsgQueue(maxSize : unsigned long)
+ send(id : unsigned long, msg* Message = NULL) : void
+ receive(id : unsigned long&) : Message*
+ ~MsgQueue()


Receiver thread use ***receive()*** function to acquire a message which has been sent to it

Item
+ id_ : unsigned long
+ msg_ : Message*

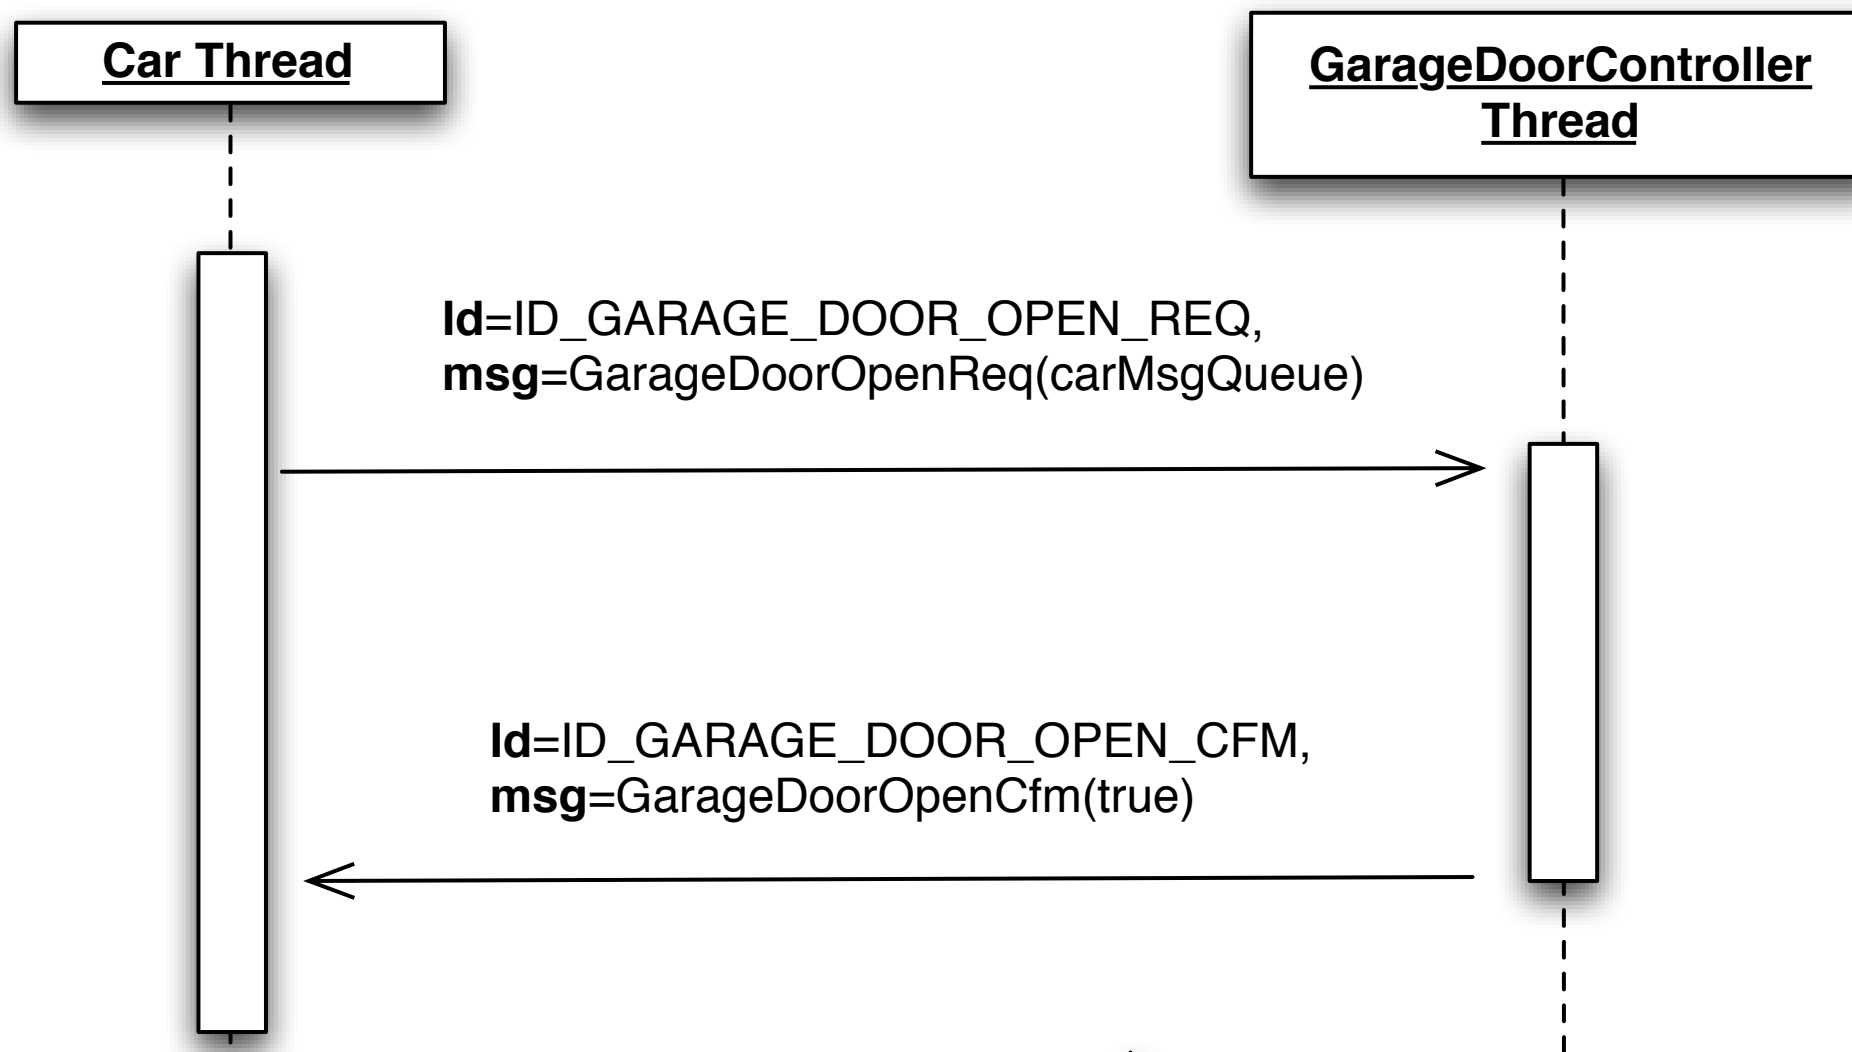
List incoming messages are placed in a queue in ***struct Item***

Case - Park-a-lot 2000

- Example: Park-a-lot 2000: An automated car parking system
 - ▶ One thread steers the car
 - ▶ Another thread steers the garage door opener



Sequence Diagram



3 types of signals

1. Request - Named XXXReq (a request requires a confirm)
2. Confirm - Named XXXCfm
3. Indication - Named XXXInd (purely oneway)

More complete example

```
void garageDoorOpenControllerHandler(Message* msg, unsigned long id)
{
    switch(id)
    {
        case ID_GARAGE_DOOR_OPEN_REQ:
            GarageDoorOpenReq* gdor = static_cast<GarageDoorOpenReq*>(msg);
            // Do stuff - call handler
            break;

        case ID_XXX:
            // ...
            break;
    }
}
```

```
void* garageDoorOpenControllerFunc(void *data)
{
    MsgQueue* mq = static_cast<MsgQueue*>(data);

    for(;;)
    {
        unsigned long id;
        Message* msg = mq->receive(id);
        garageDoorOpenControllerHandler(msg, id);
        delete msg;
    }
}
```

```
int main(int argc, char* argv[])
{
    MsgQueue garageDoorControllerMq;
    MsgQueue carMq;
    pthread_t garageDoorControllerThd;
    pthread_t carThd;
```

```
    pthread_create(& garageDoorControllerThd, NULL,
                  garageDoorOpenControllerFunc, & garageDoorControllerMq);
    pthread_create(& carThd, NULL, carFunc, & carMq);
```

```
    for(;;) sleep(100);
```

```
}
```

Park-a-lot 2000 Communication

```
class Message
{
public:
    virtual ~Message(){}
};
```

```
struct GarageDoorOpenReq :
    public Message
{
    MsgQueue* mq_;
};
```

```
struct GarageDoorOpenCfm :
    public Message
{
    bool result_;
};
```

Car Thread

```
void carSendingOpenReq()
{
    // Create request
    GarageDoorOpenReq* req = new GarageDoorOpenReq;
    req->mq_ = &carMq; // Who the requester is

    // Send it
    garageDoorControllerMq.send(ID_GARAGE_DOOR_OPEN_REQ, req);
}
```

GDC Thread

```
void handleGarageOpenDoorReq(GarageDoorOpenReq* req)
{
    // Create responds
    GarageDoorOpenCfm* cfm = new GarageDoorOpenCfm;
    cfm->result_ = openGarageDoor(); // The door is open

    // Send responds to requester...
    req->mq_->send(ID_GARAGE_DOOR_OPEN_CFM, cfm);
}
```

Car Thread

```
void handleCarOpenDoorCfm(GarageDoorOpenCfm* cfm)
{
    // Check responds
    if(cfm->result_)
    {
        driveIntoParkingLot();
    }
}
```

Typical task structure in message-based system

```
void handler()
{
    while(running_)
    {
        // get message from message queue
        switch (on state) {
            case ST_IDLE:

                switch (on message) {
                    case ID_MSG:
                        // Handle message.
                        break;
                    default:
                        break;
                }

                break;
            default:
                break;
        }
    }
}
```

Perform setup here that does not belong in constructor

Get a message, e.g. msgQueue->receive()


Switch on state, each state has separate handlers

Handle each "type" of message separately

Signal design error: Thread received something it did not expect

Execution should never reach this point

Example of a State Machine



- Checkout UML Statechart at http://en.wikipedia.org/wiki/UML_state_machine

Consequences

- Negative
 - ▶ No silver bullet by far.
 - ▶ In a performance perspective not necessarily the best solution.
 - ▶ Mostly to do with a-synchronicity, meaning that you are not guaranteed an answer but have to have some form of timeout.
- Positive
 - ▶ Does not inhibit misuse, but signifies a route that makes it “more” clear, as to what is to happen when.
 - ▶ Reduces the need for critical sections e.g. mutexes and semaphores.
 - ▶ Not blocked on a conditional/mutex while waiting

Summary

- What is it we in fact have done?
 - ▶ Entered the Event Driven Programming (EDP) paradigm
- What is EDP?
 - ▶ Reaction based programming
 - ▶ Interrupts from sensors, key input, controller directives etc.
 - ▶ Multiple *correct* paths through the code
 - ▶ For more complex code structure where the code is *not* stateless state machines are the solution - *Finite State Machine (FSM)*